

LOGGING, VISUALIZATION, AND ANALYSIS OF NETWORK AND POWER
DATA OF IOT DEVICES

A Thesis

presented to

the Faculty of California Polytechnic State University,

San Luis Obispo

In Partial Fulfillment

of the Requirements for the Degree

Master of Science in Electrical Engineering

by

Neal Nguyen

December 2018

© 2018
Neal Nguyen
ALL RIGHTS RESERVED

COMMITTEE MEMBERSHIP

TITLE: Logging, Visualization, and Analysis of Network and Power Data of IoT Devices

AUTHOR: Neal Nguyen

DATE SUBMITTED: December 2018

COMMITTEE CHAIR: Andrew Danowitz, Ph.D.
Assistant Professor of Electrical Engineering

COMMITTEE MEMBER: Bruce Debruhl, Ph.D.
Assistant Professor of Computer Science

COMMITTEE MEMBER: Lynne Slivovsky, Ph.D.
Professor of Electrical Engineering

ABSTRACT

Logging, Visualization, and Analysis of Network and Power Data of IoT Devices

Neal Nguyen

It is estimated that there are 23.14 billion IoT(Internet of Things) devices currently in use worldwide. This number is projected to grow to over 75 billion by 2025. Despite their ubiquity little is known about the security and privacy implications of IoT devices. Several large scale attacks against IoT devices have already been recorded.

To help address this knowledge gap, we have collected a year's worth of network traffic and power data from 16 common IoT devices. From this data, we show that we can identify different smart speakers, like the Echo Dot, from analyzing one minute of power data on a shared power line.

ACKNOWLEDGMENTS

Thanks to:

- My parents for supporting me throughout life. I wouldn't have made it here without your time and effort.
- Ryan Frawley for setting up the scripts for logging all of the devices as well as extra analysis on the devices alongside me. You set a really good base for us to build upon!
- Cal Poly and Cisco for funding this project. We would not have been able to get all of the IoT devices without your help.

TABLE OF CONTENTS

	Page
LIST OF TABLES	viii
LIST OF FIGURES	ix
LIST OF CODE LISTINGS	x
CHAPTER	
1 Introduction	1
1.1 Previous Work	2
1.1.1 An Analysis of Home IoT Network Traffic and Behaviour	2
1.1.2 ProfilIoT	2
1.1.3 Logging and Analysis of Internet of Things (IoT) Device Network Traffic and Power Consumption	3
1.2 Scope	3
1.3 Thesis Organization	4
2 Method	6
2.1 IoT Testbed Overview	6
2.2 Physical Layout and Setup	8
2.2.1 Wireless Access Point	8
2.2.2 Devices	9
2.2.3 Smart Power Plugs	12
2.2.4 Wiring and Configuration	12
2.3 Software	14
2.3.1 Wireless AP Code	14
2.3.2 sniff.py	14
2.3.3 Database	16
2.3.4 Device Inventory	22
2.3.5 Usage Flow	22
2.3.6 realtimeIoTGrapher.py	26
3 Results	34
3.1 Swapping the Power Switch	34
3.2 Whole Database	35

3.3	Smart Speakers	36
3.3.1	Google Home Mini	36
3.3.2	Echo Dot	38
3.4	Streaming Devices	38
3.4.1	Google Chrome Cast	39
3.4.2	Amazon Fire Stick	41
3.4.3	Roku Express	42
3.5	Summed Power Graph	43
3.6	Summed Power Graph With Noise	48
3.7	Smart Speaker Comparison	53
4	Discussion	55
4.1	Switching Smart Plug	55
4.2	Metadata Analysis	55
4.3	Power Spike at 18:12	55
4.4	Device Fingerprinting	56
4.5	Smart Speaker Comparison	57
5	Conclusion	58
5.1	Future Work	58
BIBLIOGRAPHY	60
A	How to install and run RealTimeIoTGraph.py	63
A.1	How To Install and Run	63
A.2	Usage Directions	63

LIST OF TABLES

Table	Page
2.1 IoT Devices Being Monitored	9
2.2 Columns in Network Traffic Table	17
2.3 Columns in Power Table	19
2.4 Device inventory excerpt. Password column not shown.	22
2.5 Event Log Excerpt	23

LIST OF FIGURES

Figure	Page
2.1 IoT testbed.	7
2.2 Smart Speakers Connected to WeMo Insight Switches	7
2.3 IoT Devices Under Examination	8
2.4 Navicat IP query with rollup.	18
2.5 Power query from database with Navicat.	20
2.6 Network query from database with Navicat.	21
2.7 Manual Analysis of IoT devices in Google Sheets: Raw data shown in columns A and B are then formatted into columns C and D for Graphing	27
2.8 Resulting graph of dataset shown in Figure 2.7	27
2.9 Summed power traces without interpolation.	30
2.10 Summed power traces with interpolation.	30
 3.1 Switching the Echo Dot 1 and Eufy Genie 1's smart switch	35
3.2 Internet and Transport Layer Protocols in Database	36
3.3 Idle Traffic of Google Home Over 1 Hour Period	37
3.4 Home Mini Response to News and Weather	37
3.5 Echo Dot Response to Lights	38
3.6 Chromecast First Time Boot Network Traffic and Power Consumption	39
3.7 Chromecast Video Streaming	40
3.8 Chromecast Idle Traffic	40
3.9 Fire TV Stick First Time Boot Network Traffic and Power Consumption	41
3.10 Fire TV Stick Video Streaming	42
3.11 Roku Express Video Streaming	43
3.12 5 Smart Speakers Power Summed Up. Toggled the mute button and queried each device for the weather.	44
3.13 5 Smart Speakers Power Summed Up. Queried each device for the news.	45
3.14 5 Smart Speakers Power Summed Up. Queried each device for the best basketball player.	46

3.15	5 Smart Speakers Power Usage over time. Queried each device for the best basketball player.	47
3.16	5 Smart Speakers Network throughput over time. Queried each device for the best basketball player.	47
3.17	Spike summary of for each query.	48
3.18	Idle Fan with figure 3.14 trace.	50
3.19	Idle microwave with figure 3.14 trace.	50
3.20	Idle refrigerator with figure 3.14 trace.	51
3.21	Idle PC (Intel NUC) with figure 3.14 trace.	51
3.22	In use Microwave with figure 3.14 trace.	52
3.23	Fridge in the middle of cooling with figure 3.14 trace.	52
3.24	PC in use with figure 3.14 trace.	52
3.25	Power usage of Echo Dot 1, Eufy Genie 1, and Google Home over time when asked “who is the best basketball player”. Same graph as figure 3.15 with Echo DOt 2 and Eufy Genie removed.	53
3.26	Power usage of Echo Dot 1, Eufy Genie 1, and Google Home over time when asked “who is the best basketball player”. Same time frame as figure 3.25	54
A.1	RealTimeIoTGrapher usage image	64

LIST OF CODE LISTINGS

2.1	Open and Read from a Socket	15
2.2	Rescan if all WeMos not found.	16
2.3	Open and Read from a Socket	24
2.4	Efficient SQL query to obtain total Network throughput at each second.	29
2.5	IP lookup table in real time IoT grapher.	33

Chapter 1

INTRODUCTION

There are 23.14 billion IoT devices in use worldwide; by 2025 that number is broadcast to increase to 75.44 billion[1]. With so many manufacturers creating IoT devices, each with differing update policies, some devices inadvertently have better support. For example, some devices may drop out of a manufacturer's update cycle, introducing privacy and security concerns. Several large-scale cybersecurity attacks such as Mirai [2] highlight flaws in IoT security and privacy.

To address these concerns, this thesis contributes an IoT testbed that logs network and power data from 16 IoT devices over one year, accumulating 184.94 GB of data and 172,445,929 rows into a database. To help researchers sort and view this data, this thesis adds a Python program that graphs network traffic and power data over time from the database. The graphs created by this tool were used to analyze IoT devices network and power usage in the testbed while idle, during startup, and while in use. From these graphs, it is possible to identify the smart speaker in use when viewing just one minute of the shared power usage.

This paper focuses on addressing security and privacy flaws, that if fixed, do not affect the core features of an IoT device. For example, it is required for a smart speaker to listen to its user and store the audio in increments so that it can parse the audio for the wake word. If the smart speaker occasionally hears a false positive wake word and sends the audio to its server, that is reasonable. However, Google Homes had an issue with their wake button caused the Home to listen 24/7 [3]. In another case, one of the largest IoT cybersecurity attacks, Mirai, was able to use weak login credentials to take control of 2.5 million IoT devices to perform a denial of service attack [4]. In this case, a manufacturer cannot expect a user to change their login info and should create a more secure sign-in process. This paper uses these flaws as a focus when analyzing network and power usage.

1.1 Previous Work

This section presents and analyzes related works on the topic of analyzing and characterizing IoT devices. It presents the previous works individually. Because these papers are similar to each other, commentary on how their work is different from ours and how it is useful to us as a group in section 1.2.

1.1.1 An Analysis of Home IoT Network Traffic and Behaviour

The scope and work of *An Analysis of Home IoT Network Traffic and Behaviour* [5] are the most like the work and goals of this paper. In *An Analysis of Home IoT Network Traffic and Behaviour* [5], the authors analyze IoT traffic in the home. The authors created an IoT testbed by setting up multiple IoT devices, connecting them to a router, sniffing their network packets while idle, and storing these packets on a Linux box's disk. The IoT testbed consists of a smart air quality monitor, Amazon Echo, a few Apple devices, a smart hub, and a smart vacuum cleaner.

After 22 days of network logging, the authors analyzed each IoT device individually and as a whole. For example, they noticed that they can identify most devices from the first three MAC address bits. The Hue bridge broadcasts credentials over HTTP, which is unencrypted. The authors state that these seemingly small security flaws create a privacy risk. A user's presence in a room or house can be determined from these unencrypted HTTP packets. The authors also show the percentage of network packets by protocol and various other device network patterns. This general analysis fingerprints each device.

1.1.2 ProfilIoT

The paper, *ProfilIoT: A Machine Learning Approach for IoT Device Identification Based on Network Traffic Analysis* [6] uses machine learning algorithms to classify IoT devices. The researchers of this paper collect traffic from 13 different IoT and non-IoT devices. The IOT devices include a baby monitor, motion sensor, printer, refrigerator, security camera, socket, thermostat, smartwatch, and television. The

non-IoT devices include two PCs and two smartphones for comparison. These devices connect to a Wi-Fi access point that recorded their network traffic with Wire Shark[7].

The researchers use machine learning on single-sessions to classify a device as an IoT device or non-IoT device. Then, they can classify the IoT devices by brand and model(e.g. Samsung Refrigerator, LG TV, WeMo Motion Sensor) with multi-sessions. A single-session is a 4-tuple formatted as source IP, destination IP, source port Number, destination port Number. When intercepting network traffic, they extract the information they need from each TCP packet to form the four-tuple data type. A multi-session is a list of single-sessions. Another machine learning model determines the minimum number of single-sessions needed to classify each device, determining the size of a multi-session. With single-sessions, they could determine if the device is an IoT device or not with 100 percent accuracy. Then out of their nine IoT devices, they can classify brand and model of the IoT device with 99.281 percent accuracy when run 7376 times.

1.1.3 Logging and Analysis of Internet of Things (IoT) Device Network Traffic and Power Consumption

Logging and Analysis of Internet of Things (IoT) Device Network Traffic and Power Consumption[8], written by Ryan Frawley, was formed in conjunction with this paper. Frawley's paper and this paper were both directed by advisor Andrew Danowitz at Cal Poly.

Frawley's paper documents the steps necessary to construct a reliable IoT testbed capable of capturing network traffic and power data for connected devices, and analyzing these devices further. He performed GeoIP[9] lookups on each device, showing the percentage of packets originating from each country and company. He also analyzed the packets of any unencrypted data in the devices.

1.2 Scope

The first paper from subsection 1.1.1 most closely matches this paper. The authors have the same overall idea to collect network data and then use it to analyze metadata

surrounding the networks. Our work expands on this concept by contributing a portable database consisting of 10 months of data rather than 22 days of data. This paper adds more devices in our study and focuses more on device power/network usage over time rather than specific network packet information.

Then, like the second paper from subsection 1.1.2, this paper also focuses on classification of devices from data. However, instead of using machine learning techniques on network data, this paper focuses on manual analysis, looking for spikes in power usage, the height of the spike, the length of the power spike, and other graphical heuristics.

In comparison to the first paper in 1.1.1 and second paper in 1.1.2, this paper adds power usage over time to the data set. The two papers mentioned only focus on network traffic. This paper also puts a significant emphasis on creating an extensive database rather than a smaller set of data to create graphs on network and power usage over time.

This paper is a continuation of the third paper in subsection 1.1.3. Due to overlap between these two works, certain aspects of the IoT testbed setup and usage is only covered in cursory detail here. The reader is advised to access Frawley's work for full information. We both assembled the IOT test bed and interacted with the devices on a daily basis to simulate regular usage. We both also performed a preliminary analysis of the device network and power usage together.

The unique contribution of this work is its analysis of IoT device power usage and the introduction of a custom data visualization tool that attaches to the database. This paper focuses on a select few devices, analyzing their startup, idle, and in use network and power usage over time. I compare the smart speakers' network and power usage and show that it is possible to identify a smart speaker through analysis of the powerline over time.

1.3 Thesis Organization

Chapter 2 discusses my steps in setting up the IOT test bed, the analysis tool, and the logging system for interaction with devices. Chapter 2 also highlights the steps to

set up a developer environment to run the analysis tool and how to use it. Chapter 3 presents power and network traffic for smart speakers and streaming devices while idle, in use, and during startup in the form of line graphs. Chapter 3 also shows the graphs used to fingerprint the smart speakers while handling different commands and under noise. Chapter 4 discusses the data presented in Chapter 3, why a device might have higher throughput traffic and the feasibility of classifying devices within a household from a shared power line. Finally, Chapter 5 finishes with concluding thoughts and future work.

Chapter 2

METHOD

To log and analyze network and power usage for our IoT devices, we set up an IoT testbed to log this data to our database. This chapter documents the setup process for each IoT device, the logging process, and the software created to analyze the logged data.

2.1 IoT Testbed Overview

When referring to our IoT testbed, we refer to the IoT devices and any connected hardware to log the power data.

The flow diagram in figure 2.1 outlines the IoT test bed constructed for this project. IoT devices pull power from the outlets through the smart switches, which log the power usage, and the WAP queries the smart switches for that data and pushes it to an AWS (Amazon Web Services)[10] MySQL[11] database as an entry in the ‘power’ table.

The wireless net provided through the WAP intercepts network packets of devices connected to it (IoT devices in testbed), parses them into power and network entries, and pushes them to the database. Finally, for analysis, a python script called `realTimeGrapher.py` pulls from the database and graphs this data in real time or from set intervals in the past.

Section 2.2 discusses what devices we used, the setup of each device, and the physical layout of the IoT test bed. Section 2.3 goes over the software part of the IoT test bed. It covers the usage flow of this test bed and what was done to obtain results for analysis.

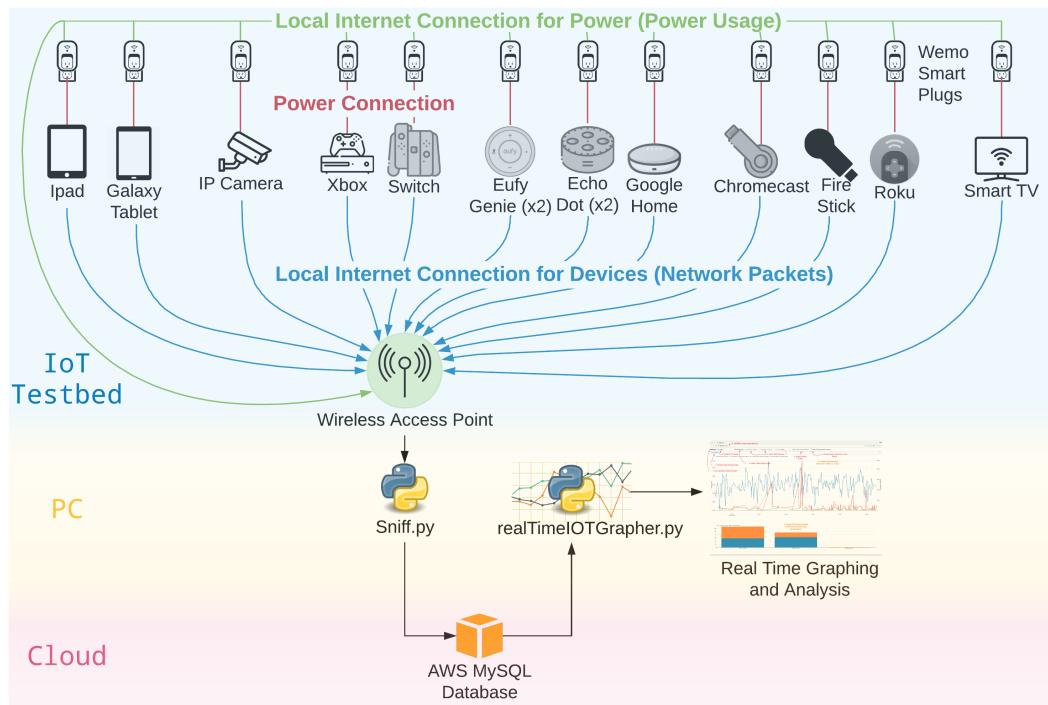


Figure 2.1: IoT testbed.



Figure 2.2: Smart Speakers Connected to WeMo Insight Switches

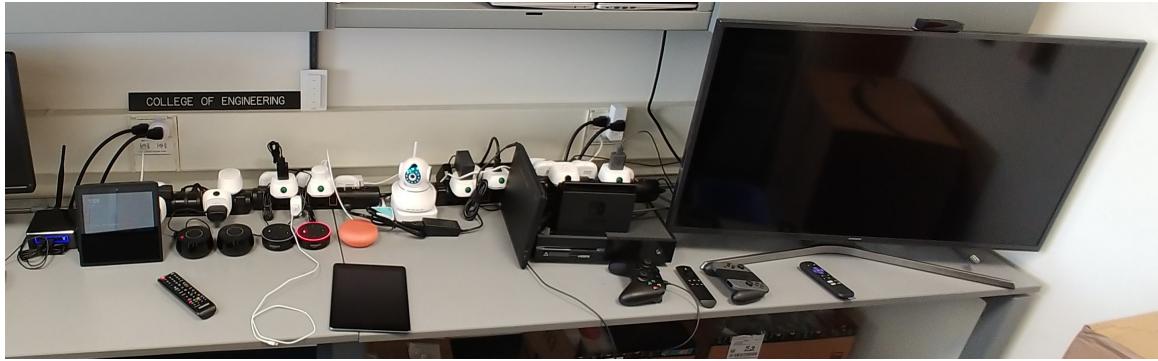


Figure 2.3: IoT Devices Under Examination

2.2 Physical Layout and Setup

This section covers the physical setup and layout of the hardware used when setting up the IoT test bed, the IoT devices we tested, and how each device is set up. Figures 2.2 and 2.3 show the IoT testbed set up. Setup was performed with Ryan Frawley and explained in his paper 1.1.3, this paper provides an updated description.

2.2.1 Wireless Access Point

The wireless access point is an Intel NUC [12]. We loaded Ubuntu[13] on it because it is the most common Linux OS [14] and most flexible for developing scripts and programs to log network packets, query power information, and push those values to a database.

One challenging thing about the NUC is that its internal wireless chip is very slow. The throughput when using the wireless chip was measured to be 12 Mb/s down on Speedtest.net by Ookla [15] when connected to with a Chromebook [16] (25 Mb/s is the minimum possible for broadband). This rate limitation caused many of the Belkin WeMo smart switches [17] to drop their connection intermittently, causing gaps in recorded power data. To solve this, we added a USB wireless antenna to the NUC. This antenna improved our internet speed to 30 Mb/s down when tested on Speedtest.net by Ookla with the same Chromebook, solving the issue where devices would lose network connectivity.

2.2.2 Devices

This section goes over the smart speakers, streaming devices, video game consoles, tablets, security cameras, or other devices as listed in figure 2.1. These devices were included based on their popularity, as the goal of the paper is to build a data set from common household IoT items.

Table 2.1: IoT Devices Being Monitored

Category	Manufacturer	Device	Quantity
Game Console	Nintendo	Switch	1
Game Console	Microsoft	Xbox One	1
Laptop	HP	Chromebook	1
Media Player	Samsung	Smart TV	1
Media Player	Google	Chromecast	1
Media Player	Amazon	Fire TV Stick	1
Media Player	Roku	Express	1
Security Camera	Eray	Hi3518 Wi-Fi Camera	1
Smart Speaker	Amazon	Echo Dot	2
Smart Speaker	Eufy/Anker	Genie	2
Smart Speaker	Amazon	Echo Show	1
Smart Speaker	Google	Home Mini	1
Tablet	Amazon	Fire 7 Tablet	2
Tablet	Samsung	Galaxy Tablet	1
Tablet	Apple	iPad	1

Smart Speakers

Smart speakers are IoT devices that combine speakers with built-in voice assistants, such as Amazon Alexa, Google Assistant, or Apple’s Siri. These devices are mainly controlled by voice commands, preceded by a wake word such as “hey Google”.

Currently, around 39 million people (16 percent of the US population) use smart speakers[18]. It is one of the top-selling IoT devices and it can act a central voice control for other IoT devices. In 2022 it is projected that 70 million US households will have at least one smart speaker (55 percent of US households) and around 175 million smart speakers total[19]. For these reasons, we included this group of devices in our research.

Within the smart speaker category, we included the Amazon Echo Dot, Google Home, and Eufy Genie. The Amazon Echo dot and the Google Home are the two leading smart speaker products. We added the Eufy Genie because it is an Amazon Alexa device that we can compare with the Echo Dot. We also included the Amazon Show, which is a voice assistant with a screen.

Streaming Devices

Streaming devices are IoT products that connect to a television or are built into a television (smart TVs) and streams videos or music from online services. In 2017 it was recorded that 70 million US households (58.7 percent of homes) had a television connected to streaming devices [20].

The specific streaming devices we look at include the Roku Express, Amazon Fire TV, and the Google Chromecast. These devices represent three popular streaming platforms (Roku, Google Chromecast, Amazon Fire TV) that have a combined user base of over 110 million users who watch content at least once a month [21].

To view content from these streaming devices, we use a Samsung smart TV. This TV also has streaming capabilities which we log but do not analyze in this paper.

Video Game Consoles

Video game consoles are systems specifically built to play video games. We use two out of the three most popular gaming devices including the Nintendo Switch which sold 17.8 million units [22] and the Xbox One, which sold 30 million units [23] as of January 2018.

Tablets

Tablets are devices that have phone operating systems running on them such as the Apple iPad, Samsung Galaxy tablet, and the Amazon Fire tablet. These devices were used for interacting with the other IoT devices, which require a smartphone or tablet for set up or control. Some of the IoT devices are limited to either Android, which runs on the Galaxy, or IOS, which runs on the iPad. Although there is no analysis on these devices in the paper, the network packets are logged. There is no way to track the power usage of these devices because they run on battery.

Security Camera

Security cameras are cameras meant to run constantly as surveillance, streaming the video footage so a user can view it from any connected device at any time. Two out of the five largest recorded cybersecurity attacks targeted security cameras [24]. The specific camera we use is the Eray camera. It is a generic Alibaba device with a weak username:password of admin:1234. A weak username and password make it very susceptible to cyber-attacks. There have also been reports that smart cameras have been found to send unencrypted data [25]. We selected this device because it has weak login credentials and is susceptible to Mirai. We wanted to see if it would get hacked by Mirai or some other attack, but we found no indication of such.

Other Devices

We also connected several other devices that are harder to categorize, analyze, or both.

We have a smart hub connected to a smart lock, an indoor room temperature sensor, and smart LED light bulbs. These devices are difficult to analyze because they communicate with the smart hub through Zigbee or Bluetooth at which point the smart hub aggregates the data and communicates with the network. To look at the network traffic caused by smart lights individually, we prevented all devices besides the smart LED bulbs from communicating with the smart hub. This isolation minimized other packet noise to the smart hub so that we can monitor the smart hub traffic as a replacement for the smart LED bulbs. However, even then, the smart hub might have extra overhead for other tasks, making isolated analysis difficult.

We also have a Chromebook which we used to control the Chromecast. However, because this is more a laptop than an IoT device, we decided to leave this out of our research

2.2.3 Smart Power Plugs

Smart power plugs are Wi-Fi capable, pass through outlets that another device plugs into. The outlets collect and transmit power usage over the network. They are critical to the power consumption logging in this research.

We used the Belkin WeMo Insight. We chose this smart plug because there are many open source Python libraries for pulling power information from them, and they are relatively affordable.

When setting these smart plugs up, we named each smart plug based on the device connected to it so that when we pull power information with our scripts, we can easily reference data to a smart plug. This informal naming led to issues when the wrong device is connected to the wrong smart plug, as discussed in section 2.3.6

2.2.4 Wiring and Configuration

We also tried to set static IPs for each device, but we decided that the average user would not set up a static IP for their device and left dynamic IPs.

The next step was to setup each device and corresponding smart plug through their corresponding setup application, filling in the device's, user name, email, network

configuration, etc. We used the device’s name as its user name, e.g. the Echo Dot was given the username “Echo Dot”.

When plugging in all the devices for the first time for power, we worked to plug each device into the corresponding WeMo smart plug and connected to our WAP during device set up so that we could instantly log power and network traffic information. It was essential to log power as soon as possible to capture a first-boot power profile for each device. Some devices had already been used, so they do not contain startup information(Xbox and the Echo Dot 1).

Set up and configuration took up a significant portion of the research time. We worked to set up each device with a separate email account. To do this, we made around 20 AOL accounts. We wanted email addresses that were not under control of the manufacturer of any of the devices. For example, we did not want to use any Gmail accounts because we did not want the Google Home or Google Chromecast to perform domain-specific optimizations. When setting up AOL accounts, we also faced a limitation of the number of email accounts tied to a single phone number. We had to use different phone numbers to circumvent this.

We set up each device in waves, taking around a week to set everything up. Incremental device setup with no plan resulted in a messy work environment, which made it difficult to match a device to its corresponding smart plug. After a month of data logging in a disorganized scheme where it was difficult to keep track of each device, we unplugged everything, renamed the WeMOS, organized the wiring, and organized the location of each smart device resulting in the set up in figures 2.2 and 2.3. With limited power strips, the organization helped maximize power ports. We also made sure to face the button for each power plug towards our view so that we could quickly notice whether the power plugs were on or off. Organization streamlined the logging process overall; debugging network and connection issues was much simpler with an ordered set up. We could quickly identify redundant devices (there are two Echo Dots and Eufy Genies), run procedures on groups of devices, and determine the corresponding smart plug for each device.

2.3 Software

This section covers the software components used to control the IoT test bed. It discusses the scripts for logging power and network traffic, the database information, and the python script for visually viewing the database.

2.3.1 Wireless AP Code

To turn the NUC into a wireless access point, we used a `create_ap` script [26]. This app takes in the Wi-Fi name, Wi-Fi password, and the wireless interface. This script turned the NUC into a wireless access point (WAP). We later found that, with newer versions of Ubuntu, there is a built-in feature for creating a hotspot through a wireless interface on the device [27]. We later used this feature to simplify the number of programs running and instead have the OS handle the hotspot. Partially through our research, as we set up more IoT devices and demand for bandwidth through the WAP increased, we had to switch to an external antenna and rerun the script to set the USB antenna as the WAP wireless interface. During this time, there is a void in network and power entries in the database.

2.3.2 sniff.py

`Sniff.py` runs a thread for power logging and a thread for network logging. As each logger runs, it writes the network packets and power information to a MySQL database hosted on AWS through the python MySQL client [28] library. If the database connection is lost, the script automatically reconnects.

Network Logging

It begins by opening a socket on the wireless interfaces connected. All IP packets are sniffed through this socket connection as shown in listing 2.1.

```

1 self.socket = socket.socket(socket.AF_PACKET, socket.
2                               SOCK_RAW, socket.htons(ETH_P_ALL))
3 self.socket.setsockopt(socket.SOL_SOCKET, socket.SO_RCVBUF,
4                       2**30)
5 self.socket.bind((self.interface_name, ETH_P_ALL))
6 while True:
7     packet, address = self.socket.recvfrom(MTU)

```

Listing 2.1: Open and Read from a Socket

Once we have sniffed a packet, we can obtain the metadata from the hex dump. We preserve this hex dump because in case we would want to extract more information from it later. From each hex dump, we extract and store the source IP, source host, destination IP, destination host, time, size, type, protocol, source port, destination port, source host, destination host, and hex dump. We explain what each field in section 2.3.3.

Power Logging

The second addition to our IoT device analysis is the power information. To obtain this information, we rely heavily on the Belkin WeMo Insights. Each WeMo connects to an outlet and each IoT device plugs into the WeMOS.

We used the PyWemo python script [29] to read power usage from the WeMOS once per second. This is the highest frequency we could poll for power. The WeMo is capable of reading both power and energy in mW and kW hours.

Once the script queries the WeMo for its power information, we relate the WeMo to the device connected to it by extracting the name of the WeMo. This information is used when pushing power data to our database.

One issue with PyWemo and the Belkin WeMOS was that they would sometimes disconnect, and the script would miss power data. To solve this issue, the script

rescans for WeMos until it finds all 16 WeMos in our testbed. The rescan code for this is shown below in listing 2.2

```
1 def scan_until_all_found():
2     print("Discovering WeMos")
3     switches = pywemo.discover_devices()
4     print("Discovered {} switches".format(len(switches)))
5     )
6     print(switches)
7     try_num = 1
8
9     while len(switches) < 16:
10         print("Did not discover enough switches, trying
11             again{}...".format(try_num))
12         switches = pywemo.discover_devices()
13         try_num += 1
14         print("Discovered {} switches".format(len(
15             switches)))
16
17     return switches
```

Listing 2.2: Rescan if all WeMos not found.

2.3.3 Database

To store all the result of the sniffed packets and power data, we push the data to a MySQL database hosted on an Amazon AWS server. As of writing, the database holds 172,445,929 entries that take up 184.94 GB of space. To interface with our database, we used a combination of the RealTimeIoTGrapher from section 2.3.6, Navicat [30], and MySQL's command line tool [31].

Network Table

The largest part of our database is the IP network table. This table contains all the network packets that have gone through the NUC WAP. It is currently 179.1 GB and contains 116,830,077 entries. It is so large because it consists of all network traffic generated over a 12-month period. In those 12 months, we did many high bandwidth tasks such as playing music or video. These entries contain the raw hex dump of the whole packet, which contains all the metadata plus data load, further contributing to the massive size of this table. The data table's rows represent a single packet with the columns shown in the table below 2.2.

Table 2.2: Columns in Network Traffic Table

Column Number	Column Name	Data Type
1	time	datetime
2	source	varchar
3	src_host	varchar
4	destination	varchar
5	dst_host	varchar
6	protocol	varchar
7	type (in/out)	varchar
8	src_port	varchar
9	dst_port	varchar
10	size	int
11	hexdump	longtext

Common SQL commands we used for our analysis include those shown in listings 2.5 and 2.6. The most useful commands generally required examining total throughput or device throughput using ROLLUP, shown in figure 2.4.

```

SELECT time, source, destination,
       CASE WHEN type IS NULL THEN 'total' ELSE type END type,
              SUM(size) AS total_throughput
FROM ip_log.ip
WHERE
    time BETWEEN '2018-05-17 16:30:00' AND '2018-05-17 16:40:00'
    AND
        (source = '192.168.12.48' OR destination = '192.168.12.48')
GROUP BY time, type WITH ROLLUP
LIMIT 20

```

time	source	destination	type	total_throughput
2018-05-17 16:30:27	192.168.12.48	184.31.3.33	outgoing	540
2018-05-17 16:30:27	192.168.12.48	184.31.3.33	total	540
2018-05-17 16:30:29	184.31.3.33	192.168.12.48	incoming	52
2018-05-17 16:30:29	184.31.3.33	192.168.12.48	total	52
2018-05-17 16:30:33	192.168.12.147	192.168.12.48	incoming	392
2018-05-17 16:30:33	192.168.12.147	192.168.12.48	total	392
2018-05-17 16:30:39	184.31.3.33	192.168.12.48	incoming	1500
2018-05-17 16:30:39	184.31.3.33	192.168.12.48	total	1500
2018-05-17 16:30:55	192.168.12.48	192.168.12.147	incoming	171
2018-05-17 16:30:55	192.168.12.48	192.168.12.147	total	171
2018-05-17 16:30:57	172.217.4.142	192.168.12.48	incoming	52
2018-05-17 16:30:57	172.217.4.142	192.168.12.48	total	52
2018-05-17 16:31:02	192.168.12.48	35.186.224.44	outgoing	94
2018-05-17 16:31:02	192.168.12.48	35.186.224.44	total	94
2018-05-17 16:31:12	172.217.4.142	192.168.12.48	incoming	855
2018-05-17 16:31:12	192.168.12.48	172.217.4.142	outgoing	569
2018-05-17 16:31:12	192.168.12.48	172.217.4.142	total	1424
2018-05-17 16:31:18	192.168.12.48	184.31.3.33	outgoing	52
2018-05-17 16:31:18	192.168.12.48	184.31.3.33	total	52
2018-05-17 16:31:35	35.186.224.44	192.168.12.48	incoming	91

Figure 2.4: Navicat IP query with rollup.

The hex dump column drastically increases the size of each entry in the network table. However, as a raw network packet, it is very flexible and can be manipulated for

many more use cases than the targeted columns we have. This flexibility is useful for future research, that may require data we did not explicitly pull out for this project

In the database, each device can be tracked down by looking for the IP address of that device in the source or destination columns. However, note that during set up, we did not set up static IPs for these devices. Because of this, a single device can be under multiple IPs within the table. We ignored this issue until a query for a particular IP stopped working. At which point we would invoke a command on the device whose IP changed to flood the database with packets from that device, query the database in the time frame of the command, and obtain the new IP. We logged all devices and their current IP in a Google Sheets [32] file covered in subsection 2.3.4.

Power Table

The power table holds the most important data for this portion of the research, and we refer to it for the majority of the paper's findings. The size of this table is much smaller than the network table. It contains 5.84 GB of data and 61,240,189 entries in the database. The significant difference in size between the IP table and power table is because each device only produces one entry power entry every second. The columns of the power table are shown below in figure 2.3.

Table 2.3: Columns in Power Table

Column Number	Column Name	Data Type
1	name	varchar
2	power_mw	int
3	time	datetime
4	today_kwh	varchar
5	on_for	varchar
6	today_on_time	varchar

When working with this database, we generally query for a range of power packets

in a given time range. We do this for either all devices, a subset of devices, or a single device. Some of the commands and results are shown below in figures 2.5 and 2.6.

```

SELECT time, name, power_mw
FROM ip_log.power
WHERE
    time BETWEEN '2018-5-26 1:30:00' AND '2018-7-26 1:40:00' AND
    (name = 'Echo Dot'
    OR name = 'Echo Dot 2'
    OR name = 'Chromecast'
    OR name = 'Roku'
    OR name = 'Samsung TV'
    OR name = 'Samsung Hub'
    OR name = 'Nintendo Switch'
    OR name = 'Xbox'
    OR name = 'Eufy Genie'
    OR name = 'Eufy Genie 2'
    OR name = 'Echo Show'
)
ORDER BY time DESC
LIMIT 20

```

time	name	power_mw
2018-07-26 01:39:59	Eufy Genie 2	1470
2018-07-26 01:39:59	Echo Dot	1340
2018-07-26 01:39:58	Chromecast	1275
2018-07-26 01:39:58	Echo Dot 2	1385
2018-07-26 01:39:58	Samsung TV	50
2018-07-26 01:39:58	Eufy Genie	1405
2018-07-26 01:39:58	Roku	2335
2018-07-26 01:39:58	Nintendo Switch	170
2018-07-26 01:39:58	Samsung Hub	2190
2018-07-26 01:39:57	Echo Show	2840
2018-07-26 01:39:57	Xbox	15785
2018-07-26 01:39:56	Eufy Genie 2	1465
2018-07-26 01:39:55	Echo Dot	1345
2018-07-26 01:39:55	Chromecast	1265
2018-07-26 01:39:55	Echo Dot 2	1375
2018-07-26 01:39:55	Samsung TV	25
2018-07-26 01:39:55	Eufy Genie	1405
2018-07-26 01:39:55	Roku	2330

Figure 2.5: Power query from database with Navicat.

time	source	src_host	destination	dst_host	protocol	type	src_port	dst_port	size
2018-06-05 05:50:00	192.168.12.77	N/A	173.194.203.188	pg-in-f188.1e100.net	TCP	outgoing	52352	443	52
2018-06-05 05:50:00	52.46.136.99	N/A	192.168.12.79	N/A	TCP	incoming	443	39864	81
2018-06-05 05:50:00	192.168.12.191	N/A	192.168.12.48	N/A	UDP	incoming	1901	52533	330
2018-06-05 05:50:00	192.168.12.191	N/A	239.255.255.250	N/A	UDP	outgoing	1025	1900	320
2018-06-05 05:50:22	216.5.2.17.202	lax17s05-in-f202.1e100.net	192.168.12.48	N/A	TCP	incoming	443	45465	52
2018-06-05 05:49:59	192.168.12.48	N/A	216.5.2.17.202	lax17s05-in-f202.1e100.net	TCP	outgoing	45465	443	52
2018-06-05 05:49:59	216.5.2.17.202	lax17s05-in-f202.1e100.net	192.168.12.48	N/A	TCP	incoming	443	45465	52
2018-06-05 05:49:59	216.5.2.17.202	lax17s05-in-f202.1e100.net	192.168.12.48	N/A	TCP	incoming	443	45465	107
2018-06-05 05:49:59	34.228.67.20	ec2-34-228-67-20.compute-1.amazonaws.com	192.168.12.101	N/A	TCP	incoming	50443	56543	40
2018-06-05 05:49:59	192.168.12.101	N/A	34.228.67.20	ec2-34-228-67-20.compute-1.amazonaws.com	TCP	outgoing	56543	50443	267
2018-06-05 05:49:59	192.168.12.101	N/A	34.228.67.20	ec2-34-228-67-20.compute-1.amazonaws.com	TCP	outgoing	56543	50443	52
2018-06-05 05:49:59	34.228.67.20	ec2-34-228-67-20.compute-1.amazonaws.com	192.168.12.101	N/A	TCP	incoming	50443	56543	60
2018-06-05 05:49:59	192.168.12.101	N/A	34.228.67.20	ec2-34-228-67-20.compute-1.amazonaws.com	TCP	outgoing	56543	50443	60
2018-06-05 05:49:59	192.168.12.101	N/A	192.168.12.101	N/A	UDP	incoming	53	52839	491
2018-06-05 05:49:59	192.168.12.101	N/A	192.168.12.1	N/A	UDP	outgoing	52839	53	66
2018-06-05 05:49:58	192.168.12.27	N/A	192.168.12.100	N/A	TCP	incoming	80	42060	40
2018-06-05 05:49:58	192.168.12.100	N/A	192.168.12.27	N/A	TCP	incoming	42060	80	40
2018-06-05 05:49:58	192.168.12.100	N/A	192.168.12.27	N/A	TCP	incoming	42060	80	40
2018-06-05 05:49:58	192.168.12.27	N/A	192.168.12.100	N/A	TCP	incoming	80	42060	40
2018-06-05 05:49:58	192.168.12.27	N/A	192.168.12.100	N/A	TCP	incoming	80	42060	1364
2018-06-05 05:49:58	172.217.11.78	lax17s34-in-f14.1e100.net	192.168.12.48	N/A	TCP	incoming	443	34176	52
2018-06-05 05:49:58	192.168.12.27	N/A	192.168.12.100	N/A	TCP	incoming	80	42060	40
2018-06-05 05:49:58	192.168.12.100	N/A	192.168.12.27	N/A	TCP	incoming	42060	80	149
2018-06-05 05:49:57	192.168.12.100	N/A	192.168.12.27	N/A	TCP	incoming	42060	80	40
2018-06-05 05:49:57	192.168.12.48	N/A	172.217.11.78	lax17s34-in-f14.1e100.net	TCP	outgoing	34176	443	52
2018-06-05 05:49:57	192.168.12.27	N/A	192.168.12.100	N/A	TCP	incoming	80	42060	52
2018-06-05 05:49:57	172.217.11.78	lax17s34-in-f14.1e100.net	192.168.12.48	N/A	TCP	incoming	443	34176	52

Figure 2.6: Network query from database with Navicat.

Within the power table, the name field is extracted from the name of the WeMo, which we manually name through the app. This can cause issues if the wrong device is connected to a WeMo. For example, the Echo Dot could have accidentally been connected to the WeMo named Google Home. There is no way to check for this besides manual examination, which we did after a few weeks, resetting all devices and causing a temporary void in power data. When we reset all devices, we also renamed some of the WeMOS. For example, we changed “echoDot” to “Echo Dot”. This way, all devices would follow the naming convention of separating words with spaces and capitalizing each word. Another possible naming convention could have been to name the WeMo by the IPs of the devices they connect to, but they are not static.

The sampling frequency for the power table is also low and missing some points. The script to query power data is not consistent and is unable to get power data at every second. To handle missing points, I interpolate them as shown in subsection 2.3.6. The interpolation process may introduce imprecise data as it fills a missing data point by extrapolating a linear path from the closest points before and after it.

2.3.4 Device Inventory

To keep track of all devices and important details about them, we log a device's service, distributor, related email, IP address, dependencies, previous usage, and password into a Google Sheets file. An example of entries into the Device Inventory are shown below in table 2.4. The service field denotes the service a device provides, for example, the Google Home is a smart speaker. The device field denotes the specific manufacturer and device name. The Email and password fields denotes the email and password used when setting up the device. Some devices such as the Nintendo Switch were setup without an account. The IP address field denotes the IP address of the device within our local network. The dependencies field denotes anything the device is connected to or uses. For example, the Roku express is connected to a WeMo for power logging and comes with a controller for use. The used flag denotes whether the device was previously used or not, at which point, we could not log first time startup data for that device.

Table 2.4: Device inventory excerpt. Password column not shown.

Service	Device	Email	IP Address	Dependencies	Used
Speaker	Amazon Echo Dot 1	amazonEchoDotSS0@aol.com	192.168.12.79	WeMo	yes
Speaker	Google Home	googleHomeMiniSS0@aol.com	192.168.12.48	WeMo	no
Streaming	Google Chromecast	googleChromecastSD0@aol.com	192.168.12.78	WeMo	no
Streaming	Roku Express	rokuExpressSD0@aol.com	192.168.12.68	WeMo, controller	no
Game Console	Nintendo Switch	n/a	192.168.12.160	WeMo, controller	no

2.3.5 Usage Flow

One of the goals of this work is to create a data set that represents the baseline of normal network traffic and power usage. To do this, we used each IoT device at least twice a week for the year length of this research. To build a proper dataset that other people could use for research, we set up a list of things that we should do for each device interfacing with it. We logged the activities into Google Sheets so that these events could be correlated to entries in the database, giving context when looking back on this data. For example, if someone were to look at the database and notice

that the power and network traffic was high for 3 minutes, they could look at our logs and see that the device was streaming music for those 3 minutes and assume that it was normal usage.

When logging events into Google Sheets, we include the start, end time, name of the device(s), action performed, and any individual notes. When naming the devices, we used the same naming as we did for the WeMos to maintain consistency from the log to the database. An example of entries into the table is shown in Table 2.5.

Table 2.5: Event Log Excerpt

Date	Start Time	End Time	Device	Event
5/25/2018	12:44:42	12:47:04	Home Mini	Ask for news
5/25/2018	12:55:07	12:56:37	Echo Dot	Ask for news
5/25/2018	12:55:57	12:56:17	Echo Show	Ask for weather
5/25/2018	12:56:44	13:03:07	Home Mini	Play music

The following subsections explains the script created to automate the event logging portion of the research. Then it explains the specific procedure we ran each group of IoT devices through whenever we interfaced with them.

Google Sheets Script

Logging information into the spreadsheet is a tedious task, especially when trying to analyze each device while running tasks on them. The code for the script is shown below in Listing 2.3.

```

1 function onEdit() {
2     var s = SpreadsheetApp.getActiveSpreadsheet().
3         getSheetByName("Event Log");
4     var r = s.getActiveCell();
5     var c = r.getColumn();
6
7     if( c == 4 || c == 5) { // checks the column
8         var dateCell = r.offset(0, -3);
9
10    if( dateCell.getValue() === '' ) { // is empty?
11        var startTimeCell = r.offset(0, -2);
12
13        // fill in the start time and date
14        var date = new Date();
15        dateCell.setValue(date);
16        startTimeCell.setValue(date.toLocaleTimeString());
17    }
18
19    if( c == 6 ) { // checks that the description is being
20        entered
21        // if so fill in the end time
22        var endTimeCell = r.offset(0, -3);
23
24        if( endTimeCell.getValue() === '' ) {
25            endTimeCell.setValue(new Date().toLocaleTimeString());
26        }
27    }
}

```

Listing 2.3: Open and Read from a Socket

Smart Speakers

Whenever interfacing with these devices, we always queried for the weather and the news. The specific phrases we used “<wake word> what’s the weather” and “<wake word> what’s the news”, we set a reminder for a random task in a random timeframe, and we muted the devices for a few minutes to see if they were still listening if we said the wake word.

Video Game Consoles

When interfacing with these devices, we played a game on the device for at least 15 minutes. Every week, we also browsed for games on the game store and downloaded free demos.

Afterward, we would turn off the Xbox and put the Nintendo Switch into sleep mode. The Nintendo Switch could not be put to sleep when docked, only when disconnected from the dock, so it was set to sleep rather than power off.

Streaming Devices

Within the Media Players section, we include the Roku Express, Chromecast, Amazon Firestick, and Samsung Smart TV. When interfacing with these devices, we watched YouTube for at least either one video or 3 minutes.

For the Chromecast, we cycled through different devices to cast videos from, including the iPad and the Chromebook. We wanted to get varied data in case the Chromecast prioritized streaming from different devices.

Tablets

In this experimental setup, we used the tablets less for investigative purposes, but more to set up and interface with other devices.

However, we still tracked the data and had a small set list of things to do on these devices. We used each tablet to browse the web, watch YouTube, and use various apps.

Security Camera

The only security camera we worked with was the Eray Security Camera.

Whenever interfacing with this device, we used the device under the app NVAS2 in the iPad for at least 2 minutes. Usage included streaming video from the camera to iPad and viewing it. We also controlled the camera through the app with pan and

tilt commands. When done experimenting with the security camera, we disconnected the tablet from the camera by using the “end stream” button in the NVAS2 app.

2.3.6 realtimeIoTGrapher.py

This section covers a visual tool I created to look for trends. We initially graphed the data into Google Sheets and graphed it. The tables and formulas are shown in figure 2.7. This process was tedious and time-consuming. We had to copy the data over after making a SQL query, extract unique time values, and then get the total network and power usage at each time interval.

To address these problems, I created a python script that automatically forms graphs, including the one shown in figure 2.10 when given the time range and the devices. This automation sped up the analysis process and provided extra features covered in section 2.3.6.

This Python script leverages the Plotly library, which, given data, graphs it onto a local web page for viewing. Plotly comes with many useful tools for further analysis and can be extended to run on a public web page for public viewing.

I plan to distribute this out with our database so that researchers can view the power and network information of any device at a glance.

	A	B	C	D	
1	time	size	time	SUMIF(A:A, C2, B:B)	
2	2018-04-17 11:25:40	331	2018-04-17 11:25:40	331	
3	2018-04-17 11:25:41	272	2018-04-17 11:25:41	2144	
4	2018-04-17 11:25:41	390	2018-04-17 11:25:42	1183	
5	2018-04-17 11:25:41	399	2018-04-17 11:25:43	2043	
6	2018-04-17 11:25:41	396	2018-04-17 11:25:44	865	
7	2018-04-17 11:25:41	334	2018-04-17 11:25:45	431	
8	2018-04-17 11:25:41	353	2018-04-17 11:25:46	7916	
9	2018-04-17 11:25:42	40	2018-04-17 11:25:47	8153	
10	2018-04-17 11:25:42	52	2018-04-17 11:25:48	3580	

Figure 2.7: Manual Analysis of IoT devices in Google Sheets: Raw data shown in columns A and B are then formatted into columns C and D for Graphing

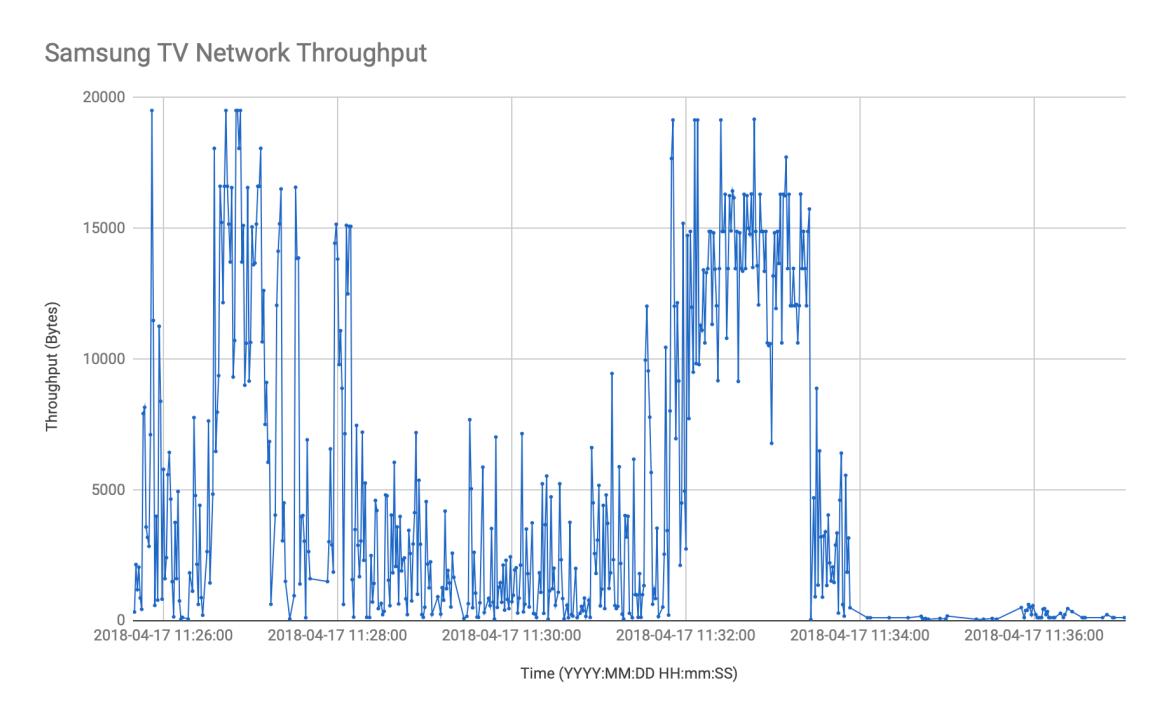


Figure 2.8: Resulting graph of dataset shown in Figure 2.7

Features

This section discusses the features that RealTimeIoTGrapher (grapher) provides.

At its core, this program automates the creation of the graph shown in figure 2.8. The user can specify a time range that the graph should cover and what devices to graph. By default, the graph shows power, input traffic, output traffic, and total network traffic over the time range specified.

The grapher also annotates the data by including a line denoting the average for each of the traces over the time range currently displayed. It also annotates the maximum and minimum values in the time range for each trace. The grapher can also display information in real time for a user settable interval amount of time, updating every second.

The Plotly libraries also provide useful tools when displaying these IoT graphs. It is possible to zoom in and out the displayed graph, select specific traces for viewing, save the graph, and hover over data points to display the specific value. Finally, once a graph is saved Plotly allows further editing of the graph so that it is formatted the way the user wants.

```

1  def throughput_query_in_range(db_connection, device,
2      start_time, end_time):
3
4      sql_query = """
5          SELECT
6              time,
7                  CASE WHEN type IS NULL THEN 'total' ELSE type
8                      END type,
9                  SUM(size) AS total_throughput
10             FROM ip_log.ip
11             WHERE
12                 time BETWEEN '%s' AND '%s' AND
13                     (source = '%s' OR destination = '%s')
14             GROUP BY time, type WITH ROLLUP;
15
16 """ % (start_time, end_time, ip[device], ip[device])
17
18     dataframe = pd.read_sql_query(sql_query, db_connection)
19
20     return dataframe

```

Listing 2.4: Efficient SQL query to obtain total Network throughput at each second.

To reduce local computation, the grapher offloads work to the server through a SQL query. When querying for network data, it performs a SQL ROLLUP to sum the incoming and outgoing network throughput to obtain the total throughput. ROLLUP is required because it provides nested queries and causes the query to sum up the data by time for each type of packet. This query is formatted as shown in listing 2.4.



Figure 2.9: Summed power traces without interpolation.

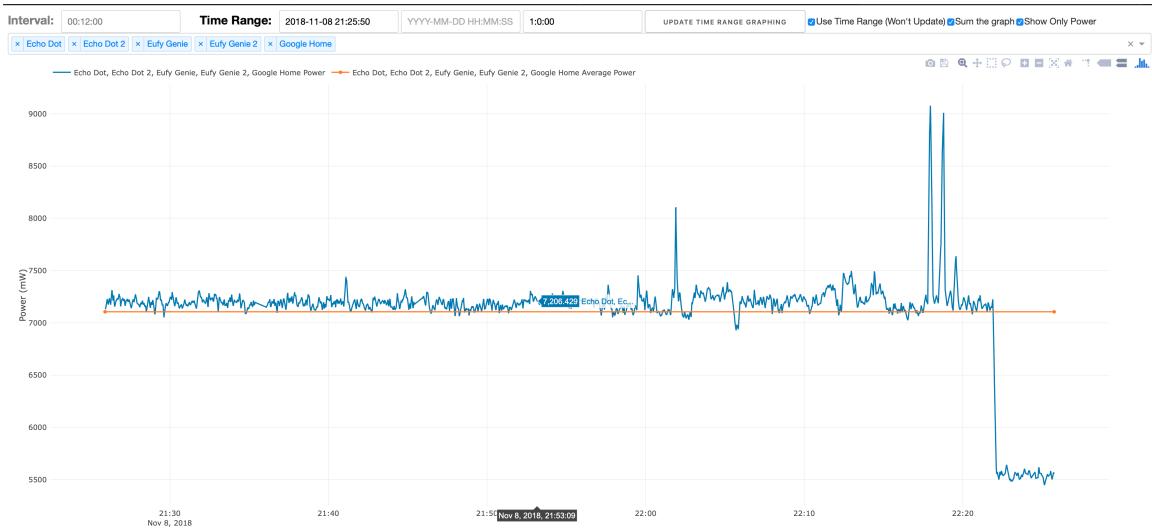


Figure 2.10: Summed power traces with interpolation.

The grapher can also sum up all power traces into one single trace. Originally, when summing the graphs, it would add the two Pandas [33] data frames that represent each trace, and graph the result, shown in figure 2.9. Pandas is a Python library that represents a multidimensional array as a table, providing useful methods for manipulating the dataset. However, the WeMo does not always sample every second. This is an issue if, at time, one device X has a power point, but another device does not. The summation counts the empty data point as 0, resulting in an erroneous sum.

This was solved by interpolating data for every second for each device before summing data frames, with a built-in Pandas feature. Interpolation is shown in 2.10 as a graph.

Visual Takeaway

This tool created most of the images shown in section 3. A few other researchers such as Frawley’s have used this tool for data visualization. The grapher provides real-time or historical graphs on any in the database device for individual analysis or comparison with little effort or time. Additional features from Plotly also provides an easy way to format graphs for presentations or papers.

The summation feature can simplify the graph and provide a more holistic view of power usage, making it easier to spot patterns. In this paper, the summation feature is used to simulate the concept of a household’s “powerline” where all energy usage would be summed up into one power meter.

Limitations

As this tool was created entirely for research purposes, there are some edge cases that it fails to handle.

When the database has too many connections, the software refuses to connect, and the grapher will not work. When graphing in real time, the grapher makes a new connection for each query. If a query fails, connections can quickly build up.

The graphing tool also slows as the time frame requested increases. We have noticed a slow down for queries with time frames longer than seven hours. In the static graphing mode, the tool slows down before graphing the information. However, in real time graphing mode, the tool makes many slow queries, eventually causing too many connections to the database and undefined behavior.

The graphing tool uses a lookup table that maps the IP address of the device to the corresponding name given to the WeMos. Because we manually create the lookup table, when any of these fields change, the graphing is not able to find the device. In some cases, this causes the graphing tool to crash. To solve this, the user must

manually change the IP addresses defined in the lookup table in lines 23-39 shown in listing 2.5 below.

To figure out the new IP address for the table, the user must perform a SQL query in the time frame that a device is doing something, and manually examine the network traffic to see which IP has the highest network throughput. With this method, we had to ensure only one device was in use so that an IP address for another device would not be confused for the current device.

```

1      #region IP Addresses of all devices we are tracking
2
3      ip = {
4          'Chromecast':           '192.168.12.77',
5          'Echo Dot 2':          '10.42.0.132',
6          'Echo Dot':            '10.42.0.150',
7          'Eufy Genie':          '10.42.0.223',
8          'Eufy Genie 2':        '10.42.0.172',
9          'Fire Stick':          '192.168.12.113',
10         'Google Home':        '10.42.0.236',
11         'IP Camera':          '192.168.12.58',
12         'Nintendo Switch':    '192.168.12.160',
13         'Roku':                '192.168.12.68',
14         'Samsung Hub':         '192.168.12.100',
15         'Samsung TV':          '192.168.12.191',
16         'Smart Light':         '192.168.12.27',
17         'Xbox':                 '192.168.12.251',
18         'Echo Show':           '192.168.12.122',
19         'Appliance':            '192.168.12.122',
20         'Appliance1':          '192.168.12.122',
}

```

Listing 2.5: IP lookup table in real time IoT grapher.

Chapter 3

RESULTS

This chapter covers different network and power usage patterns for the IoT devices. This chapter starts with the effects of changing the smart plug that the device connects to in section 3.1. Then section 3.2 covers the protocol spread over the whole database. The next two subsections 3.7 and 3.4 covers smart speakers and streaming devices while idle, in use, and during first boot, including other interesting results if any. Sections 3.5 and 3.6 covers the summed power graphs that simulate a shared powerline with and without noise. Finally this chapter concludes with a comparison between smart speakers in section 3.7.

Sections 3.2, 3.3, and 3.4 are covered in Ryan Frawley's paper [8] and are paraphrased with no further discussion in section 4. Frawley's paper contains a more in-depth analysis in those sections and include GeoIP information that highlights the spread of locations the database and each device communicates with the most.

When obtaining data, most of the packets are encrypted, so the data we present focuses on metadata such as size, protocol, source IP, dest IP. For power analysis, we only use the power used at each second.

3.1 Swapping the Power Switch

To check if the smart switches were consistent between devices we swapped various devices with various power switches. After WeMo swaps with various devices, no significant change was found when swapping the switches. One example is shown in figure 3.1 where we swapped the Echo Dot 1 and Eufy Genie 1s' corresponding Wemo smart switch. After starting restarting, the Eufy Genie's power reads the same as the Echo Dot had previously read.



Figure 3.1: Switching the Echo Dot 1 and Eufy Genie 1's smart switch

3.2 Whole Database

When examining the whole database, most of the traffic is TCP packets. The second most significant protocol is UDP. The spread of protocols can be shown in figure 3.2.




Figure 3.2: Internet and Transport Layer Protocols in Database

3.3 Smart Speakers

This section will cover the smart speaker's power and network usage in different scenarios.

3.3.1 Google Home Mini

The Google Home Mini's idle traffic broadcasts SSDP (Simple Service Discovery Protocol) packets once every minute. The SSDP packets are a discovery request packet for every IoT device on the network that supports UPnP (Universal Plug and Play). At which point devices such as the Echo Dot, Samsung TV, streaming devices, and the Chromebook respond with information about themselves in a .xml file. This XML file contains details regarding the device operating system and more. Google also sends encrypted TCP packets to Google every 10 minutes. The Google Home Mini's idle traffic is shown in figure 3.3. The high peaks are the encrypted TCP packets, and the smaller spikes are the SSDP/UPnP packets.

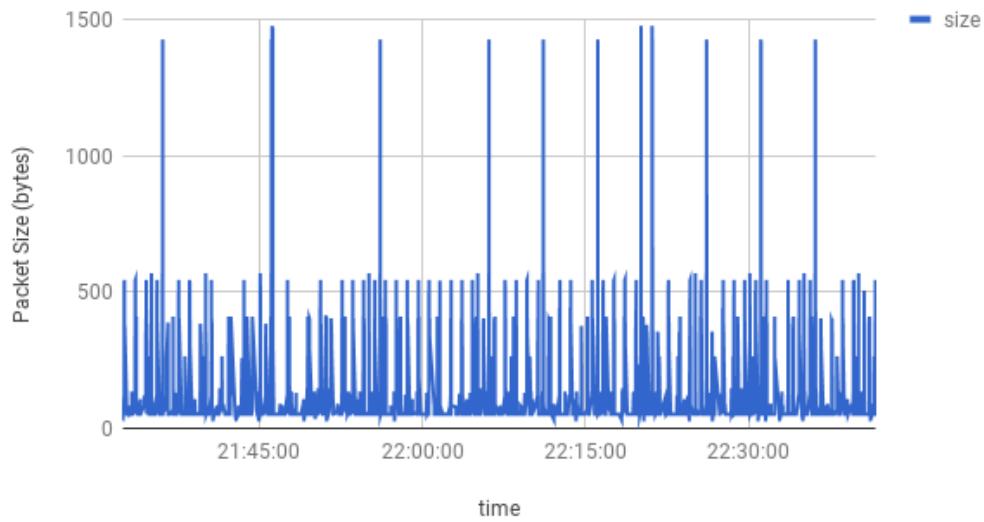


Figure 3.3: Idle Traffic of Google Home Over 1 Hour Period

A graph of the Google Home Mini is shown in the figure 3.4. The Google Home Mini is first asked for the news and then told to stop. After waiting for 40 minutes, we ask it for the weather forecast.




Figure 3.4: Home Mini Response to News and Weather

3.3.2 Echo Dot

One interesting finding from the Echo Dot is that it has a light monitor. When turning on the lights, the LEDs on the Echo Dot brighten to adapt. The brightness change causes the Echo Dot to use more idle power as shown in figure 3.5.

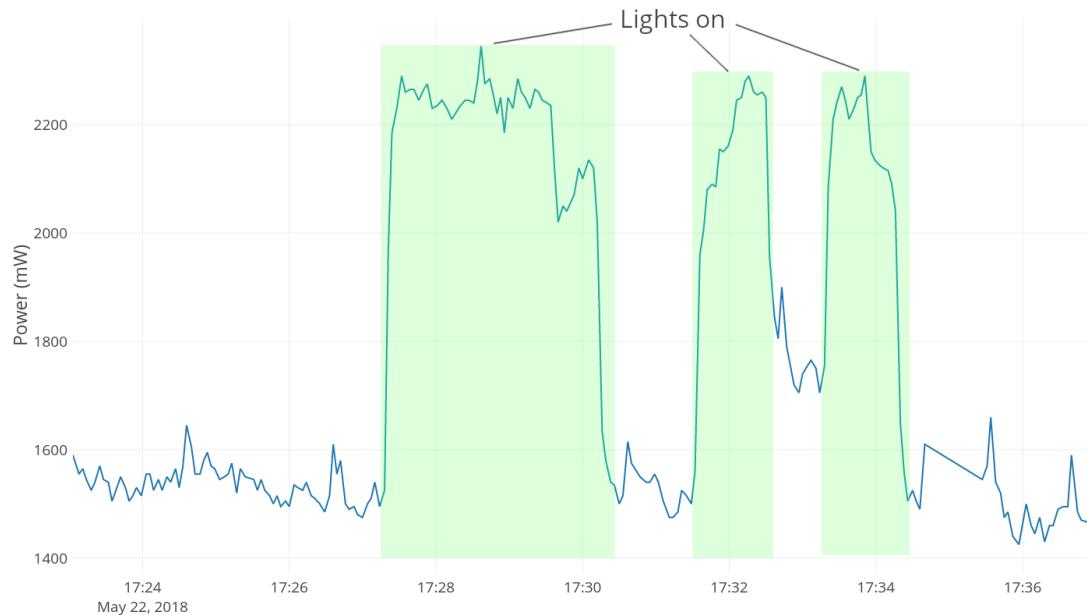


Figure 3.5: Echo Dot Response to Lights

3.4 Streaming Devices

Each subsection shows the streaming devices usage during startup and when streaming.

3.4.1 Google Chrome Cast

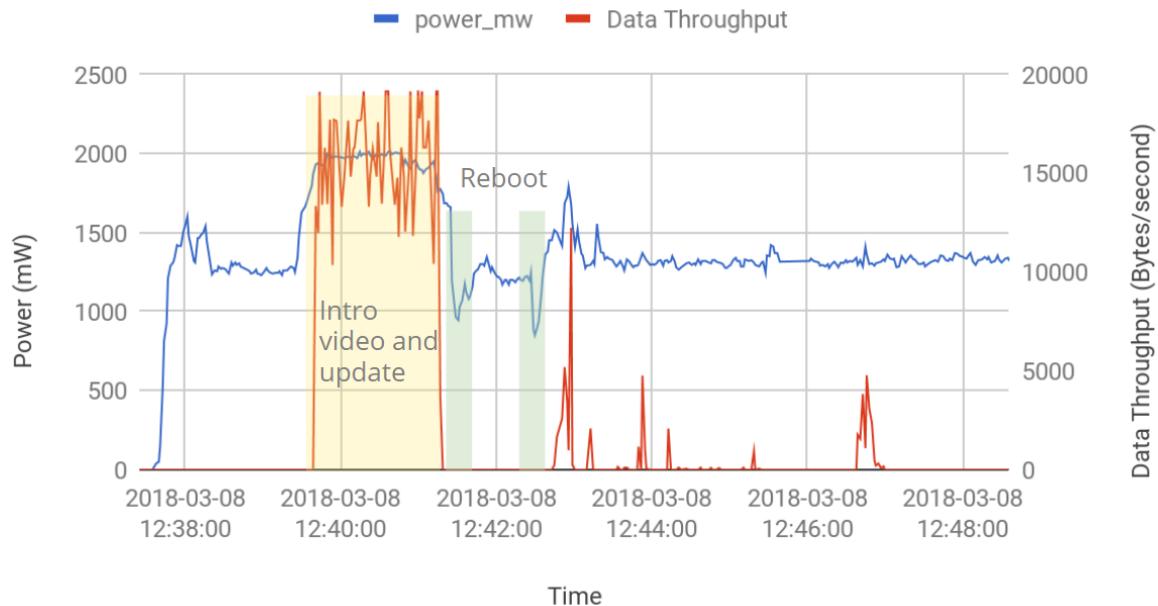


Figure 3.6: Chromecast First Time Boot Network Traffic and Power Consumption

The Chromecast startup graph is shown in figure 3.6. On startup, the Chromecast shows an intro tutorial video and downloads a firmware update. After rebooting twice, it is ready to go.

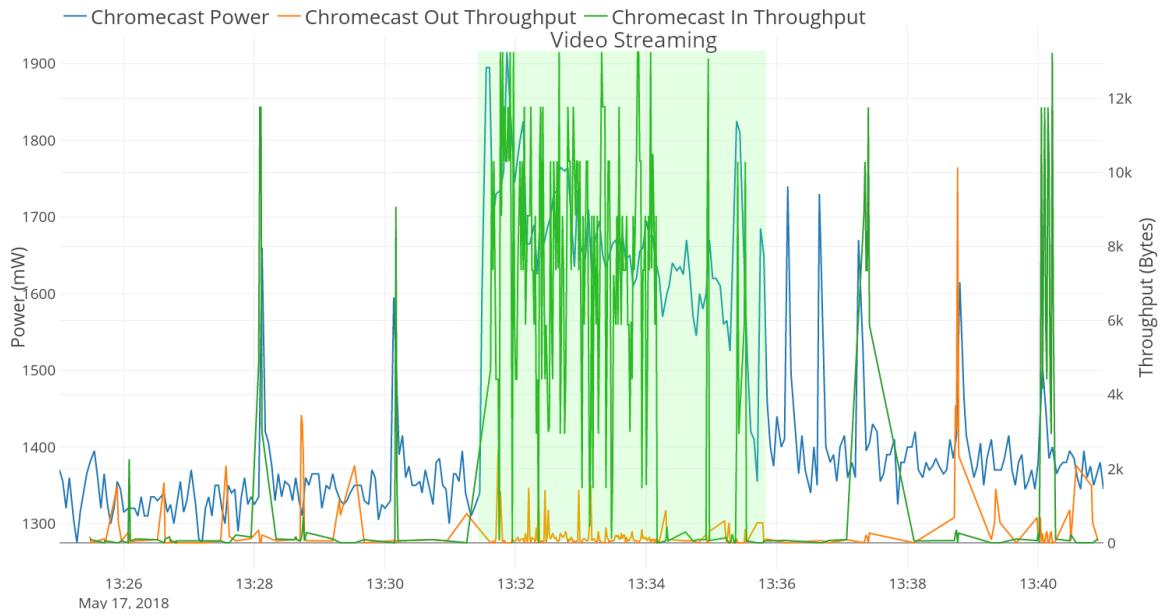


Figure 3.7: Chromecast Video Streaming

The Chrome Cast streaming graph is shown in figure 3.7. In this time frame, the chrome cast streams video in the marked box.




Figure 3.8: Chromecast Idle Traffic

Additionally, the Chrome Cast idle graph is shown in figure 3.8. In this time frame,

there are consistent spikes to the chrome cast every 2 minutes. During these spikes, the chrome cast is showing a new background that it downloads from Google Servers.

3.4.2 Amazon Fire Stick

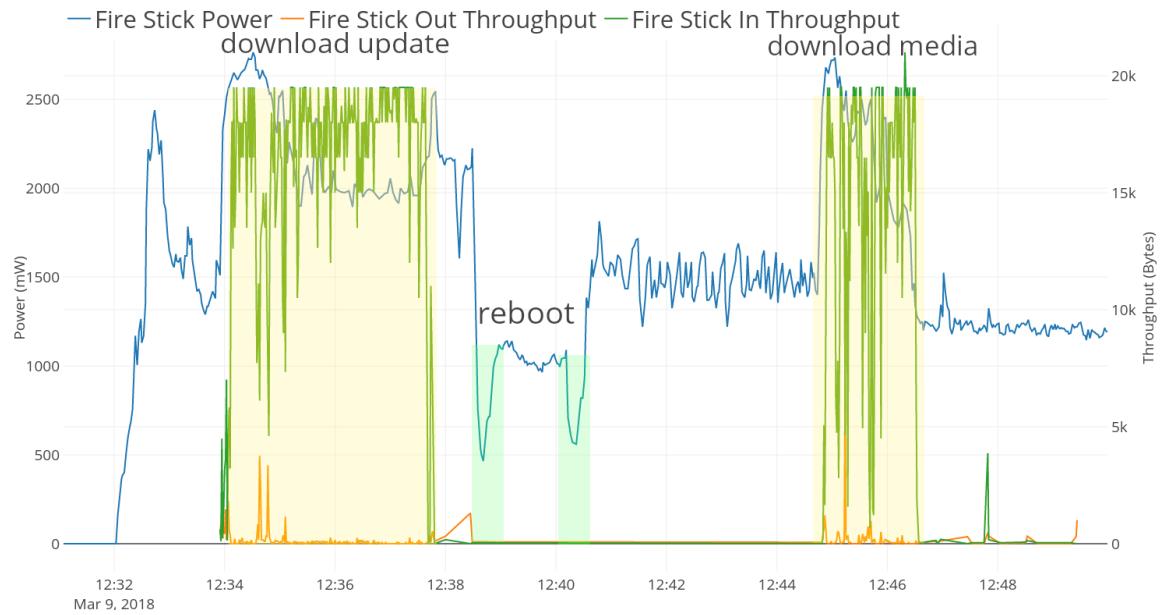


Figure 3.9: Fire TV Stick First Time Boot Network Traffic and Power Consumption

The Amazon Fire Stick startup graph is shown in figure 3.9. On first boot, the fire stick downloads an update, reboots twice, then downloads certificates from Symantec and Verisign.

When streaming, the Fire Stick increases throughput and power usage with a spike at the beginning and end of streaming in power usage as shown in figure 3.10.

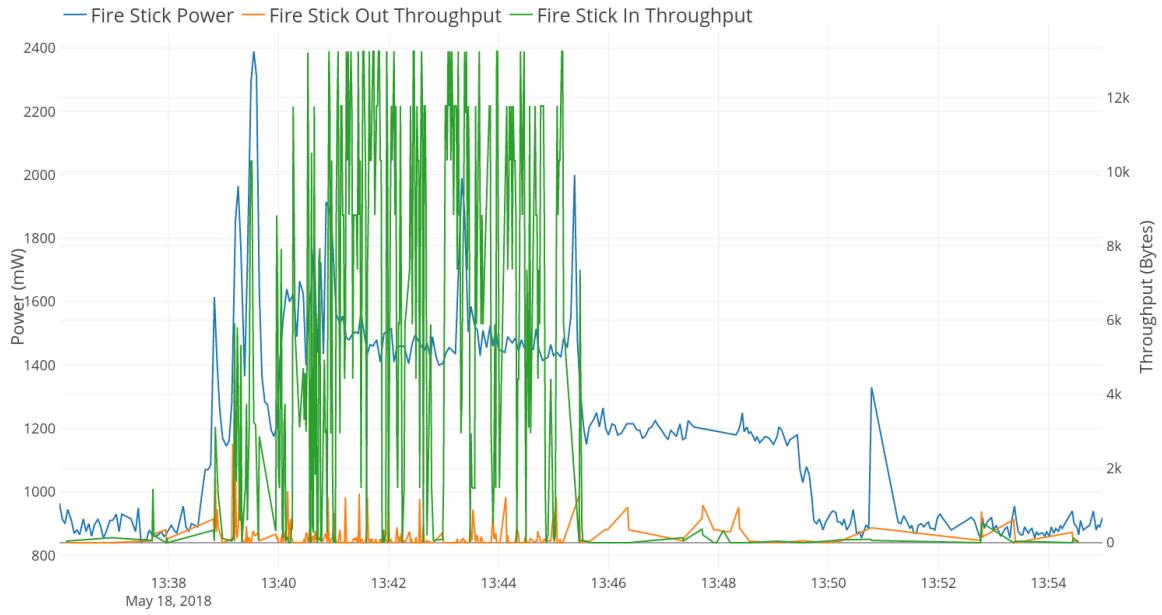


Figure 3.10: Fire TV Stick Video Streaming

3.4.3 Roku Express

When streaming, the Roku's power usage and network throughput rise and stay at a constant level until the video streaming is complete. At which point it drops back down when done.



Figure 3.11: Roku Express Video Streaming

3.5 Summed Power Graph

In this section, we show graphs of the five smart speakers summed up as one graph of power usage. The devices used includes two echo dots, 2 Eufy genies, and 1 Google chrome cast.




Figure 3.12: 5 Smart Speakers Power Summed Up. Toggled the mute button and queried each device for the weather.

The first graph in figure 3.12 is shown in a time frame where each device was unmuted and muted for a period. This shows if muting and unmuting the devices affected anything. In this graph, there are changes on the first runthrough, but when repeating the process nothing happens. Further runthroughs show that the power usage stayed the same. In the last half of the graph we queried each smart speaker for the weather while all other smart speakers were muted at the 18:30 mark. We queried each device in order 3-4 times in consecutive order so that it is easy to see and to check for consistent results. The Eufy had the smallest spike for this command at 400 mW, then the google home at 600 mW, and finally the echo dot at 2000 mW.




Figure 3.13: 5 Smart Speakers Power Summed Up. Queried each device for the news.

The next graph 3.13 shows the smart speakers power usage when asked for the news. Each speaker was asked for the news at random as to prevent any optimizations. This graph and the rest of the summed graphs in this section use the same notation for signifying commands. Two corresponding dots of the same color signify the start and end of the command for a specific device. If the dots are on another level, then this command has been repeated another time.

In this graph, all devices have a spike at the beginning and end of the command and maintain a steady energy usage in between that is slightly higher than the idle energy used. The peak to peak spike of the Eufy Genie is the smallest at 350 mW, then the Google Home at 500 mW, and finally the Echo Dot at 1900 mW.




Figure 3.14: 5 Smart Speakers Power Summed Up. Queried each device for the best basketball player.

The next graph 3.14 shows the smart speakers when asked for the best basketball player. The annotation scheme is the same as before. Each smart speaker is queried for the best basketball player in consecutive order four times. Similar to before, the Eufy has a power spike of 420 mW peak to peak, the Google Home has a power Spike of 720 mW, and the Echo Dot has a power spike of 2180 mW.

However, when looking at graph 3.14, there is a power spike that is unaccounted for in correspondence to the event log. We made a query, invoking all other power spikes but we did not do anything for the power spike at 18:12 that is higher than the other Eufy power spikes.

To figure what the power spike at 18:12 is, we separated the graphs into individual power traces as shown in figure 3.15. From the graph, the power spike at 18:12 is attributed to the Echo Dot 2. We then looked at the individual network usage for each of these devices in this time frame as shown in figure 3.16. At 18:12, there is no significant network usage. We had no conclusive evidence to decide what the power spike is accounted to, but section 4 covers some speculated reasons.

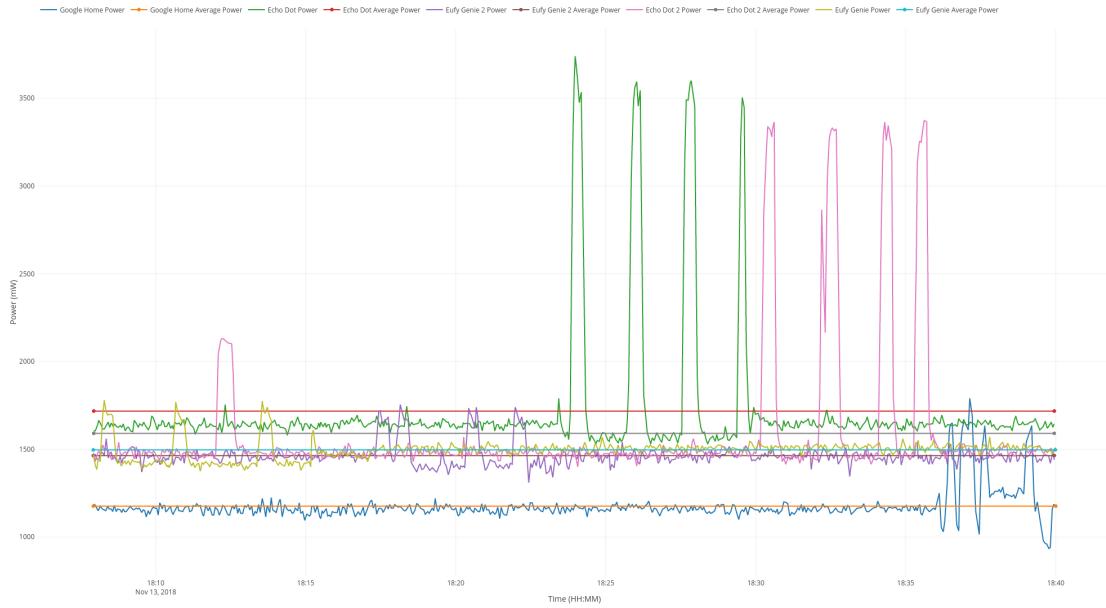


Figure 3.15: 5 Smart Speakers Power Usage over time. Queried each device for the best basketball player.

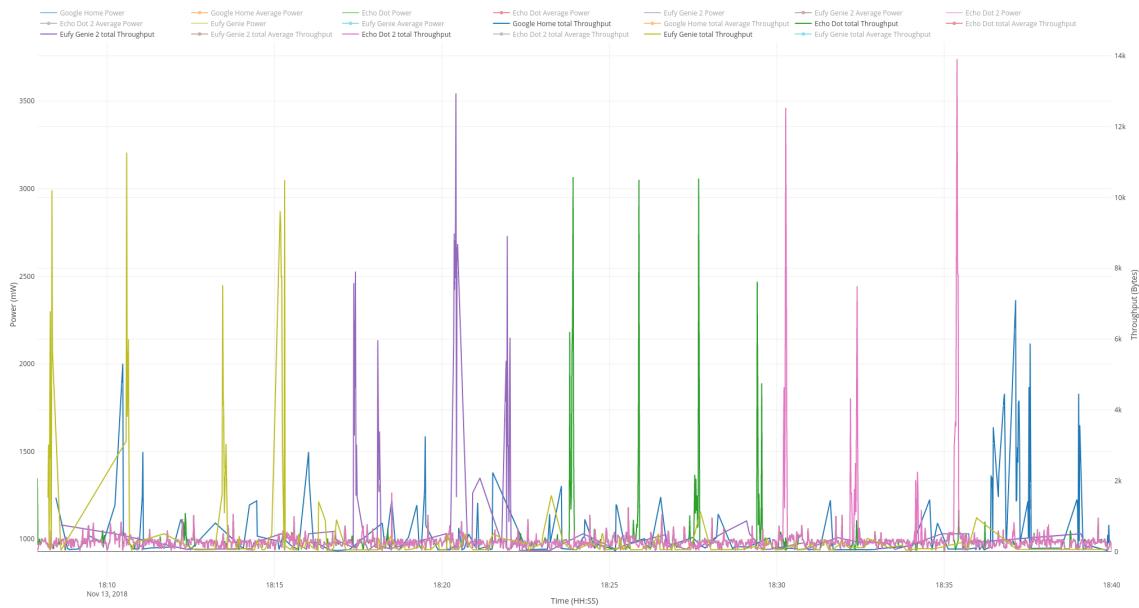


Figure 3.16: 5 Smart Speakers Network throughput over time. Queried each device for the best basketball player.

The power spike information for each query type (“what’s the weather” “what’s the news”, “who’s the best basketball player”) are shown below in figure 3.17. This bar graph also includes the average spike height for each device over these three commands.

The black line at the top of each bar shows the standard deviation. Averaging peak to peak voltage spikes for each device shows the Eufy Genie with a 390 mW spike with 36.1 mW standard deviation, the Google Home with a 606.7 mW spike with 110 mW standard deviation, the Echo Dot with a 2026.7 mW spike with 149.89 mW standard deviation.

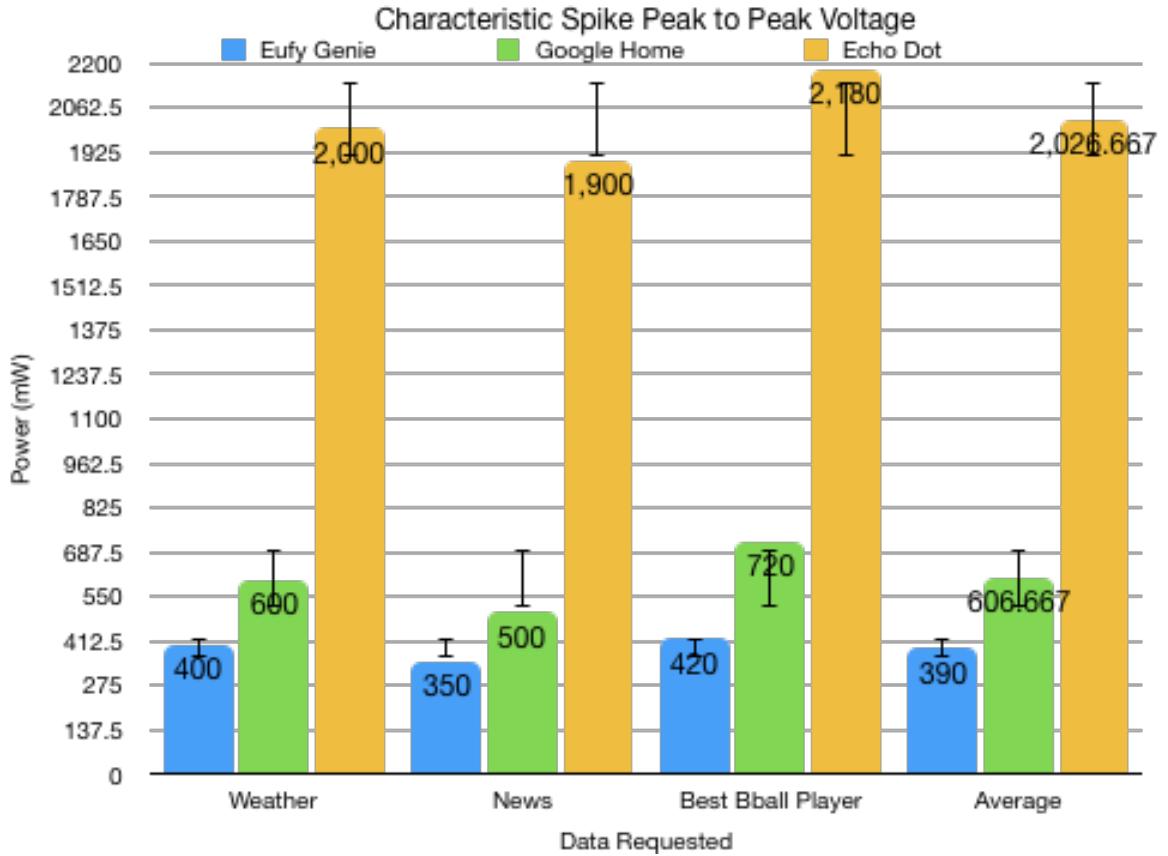


Figure 3.17: Spike summary of for each query.

3.6 Summed Power Graph With Noise

In this section, we introduce noise into the system so that we can determine if the power spike from each smart speaker is still discernible within a summed power graph when we add high power devices.

In each subsection, there are three traces. The first trace (blue trace) is the power usage over time for a high power device. This trace maps to various devices as we switch them out. The second trace (green trace) is the summed power usage of our five

smart speakers (2 Echo Dots, 2 Eufy Genies, 1 Google Home). For this second trace, we used the same power usage graph as figure 3.14 for every graph in the subsections below. Using the same graph provides our control for each smart speakers power spike. The third trace (yellow trace) is the sum of both trace 1 and 2. It is the power usage of the smart speakers in the green trace with the noise from the device in the blue trace.

In all these graphs there are be two y-axes which represent the power (mW) used by a device at that time. The graphs have two y-axes so that the traces can be more zoomed in. Usually, one trace is significantly larger than the other, so if they are on the same scale, the other gets compressed. The trace for the high power device (blue) is always in reference to the first y-axis on the left side. The trace for the smart speakers (green) is always in reference to the second y-axis on the right side. The sum of the two traces (yellow trace) is in reference to the y-axis that it is closer into magnitude. Note that to further zoom in on the graphs, the y-axes do not start at 0.

The x-axes in these graphs are also not actual time stamps. The x-axis is the time elapsed (s) from when the green smart speaker trace (green) 3.14 began. Because we had already recorded the smart speaker trace (green), we recorded each high power appliance trace (blue) individually and included it on the same time range.

Below, we display the smart speaker trace with a PC (Intel NUC), fan, refrigerator, or microwave. We do this while these devices are idle and in use.

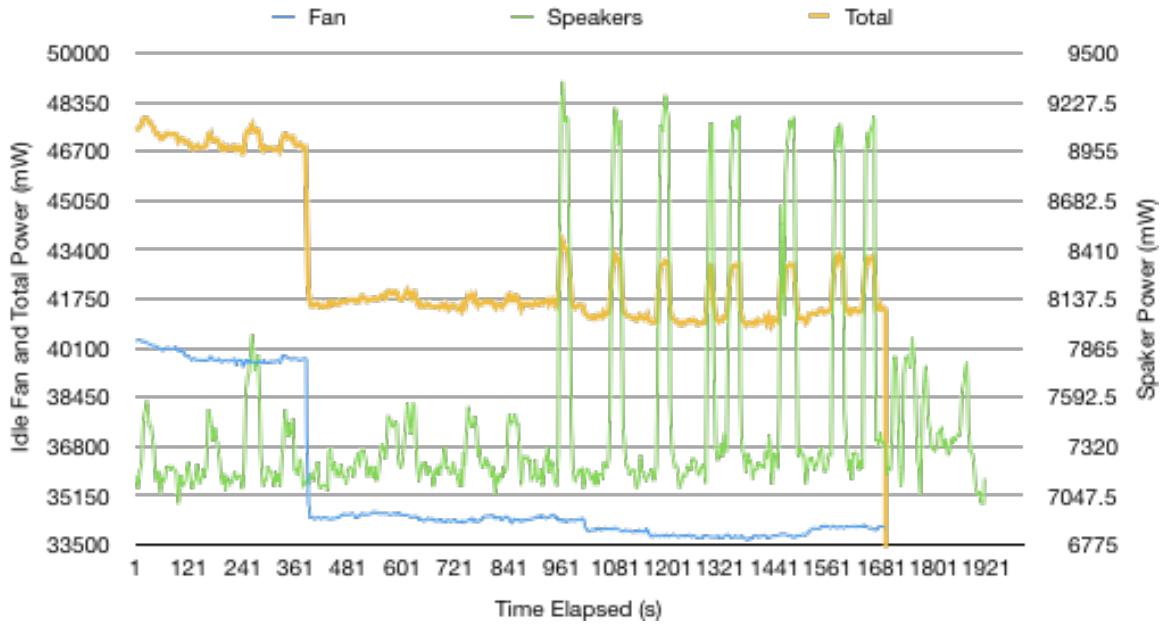


Figure 3.18: Idle Fan with figure 3.14 trace.




Figure 3.19: Idle microwave with figure 3.14 trace.

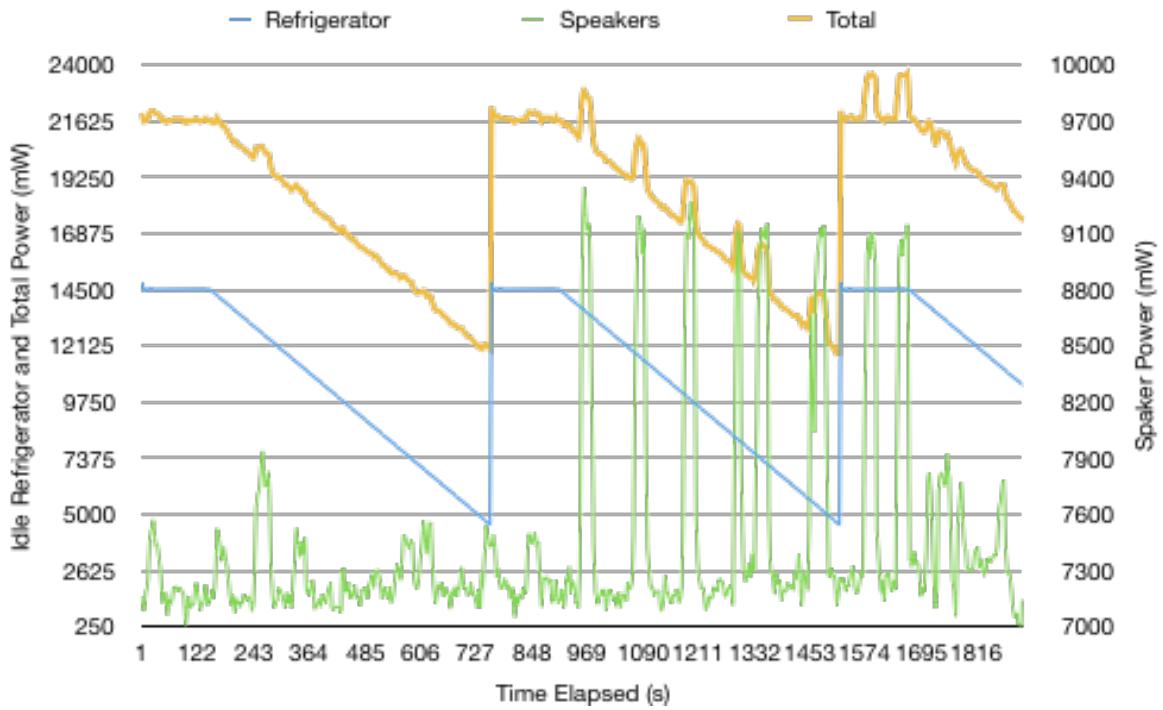


Figure 3.20: Idle refrigerator with figure 3.14 trace.

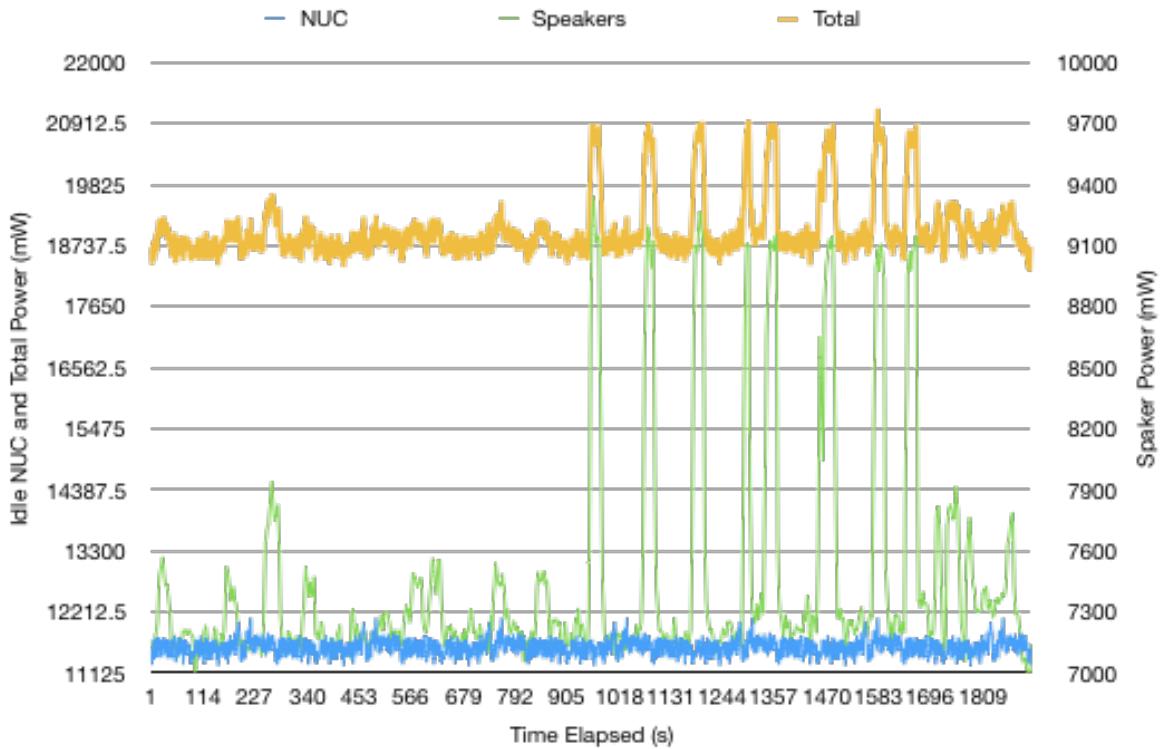


Figure 3.21: Idle PC (Intel NUC) with figure 3.14 trace.

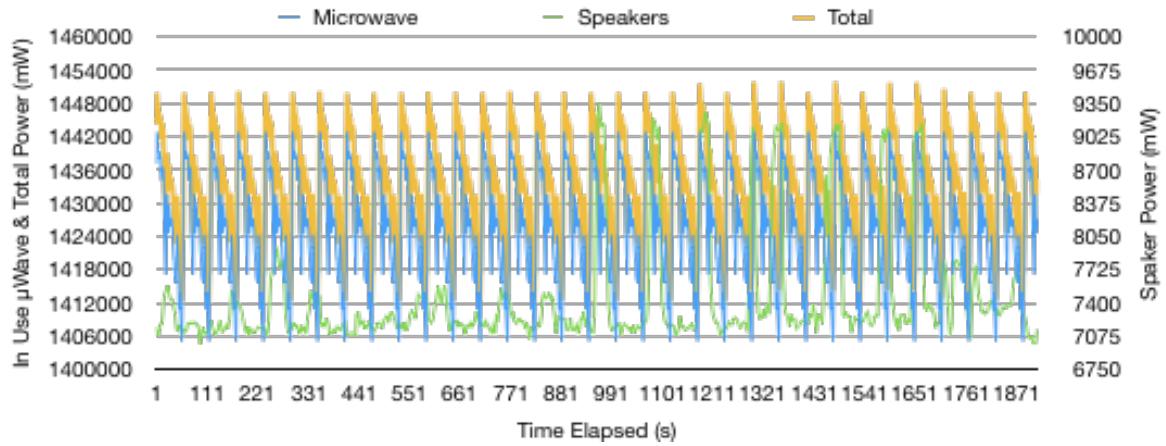


Figure 3.22: In use Microwave with figure 3.14 trace.

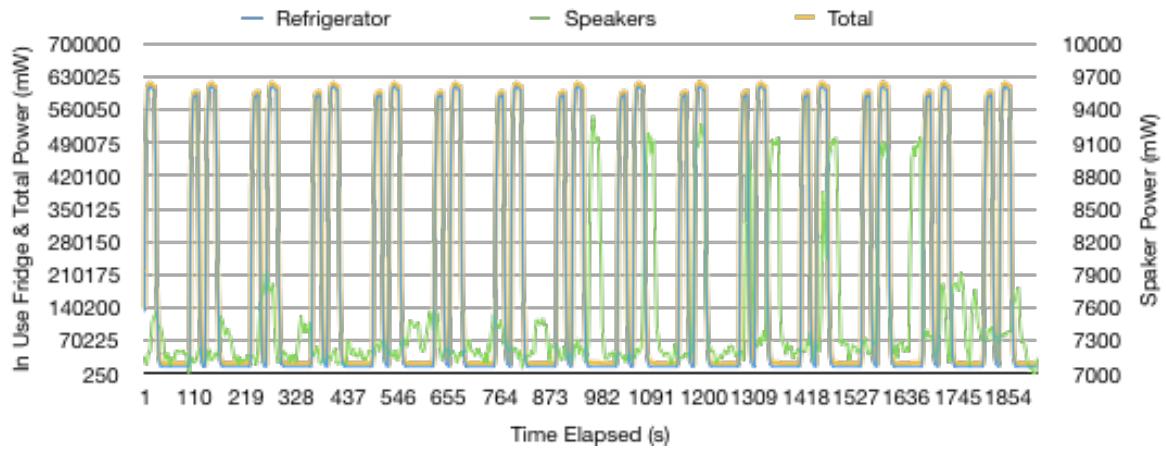


Figure 3.23: Fridge in the middle of cooling with figure 3.14 trace.

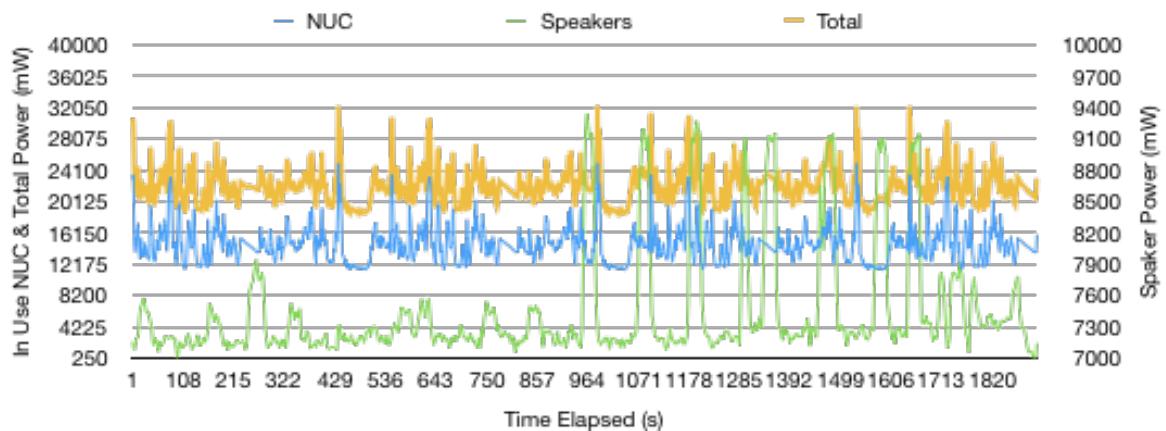


Figure 3.24: PC in use with figure 3.14 trace.

3.7 Smart Speaker Comparison

The figures in section 3.5 show that the Echo Dot uses a lot more energy than the other smart speakers. Because of this, we wanted to compare the energy and network usages of the three smart speakers individually so that we can determine trade-offs that they make.

In the graphs below, we display power and network traces for the echo dot 1, Eufy genie 1, and Google home. We used the same time frame as the graph in figure 3.15 for the two graphs below. We removed the second Echo Dot and Eufy Genie for simplicity.

In the same time frame as the two graphs below, we also recorded the average power usage and throughput for the smart speakers. For average power, from greatest to least, we have the Echo Dot (1650 mW), Eufy Genie (1325 mW), Google Home (1175 mW). For average throughput, from greatest to least, we have the Eufy Genie (1743 bytes), Google Home (800 bytes), Echo Dot (325 bytes).

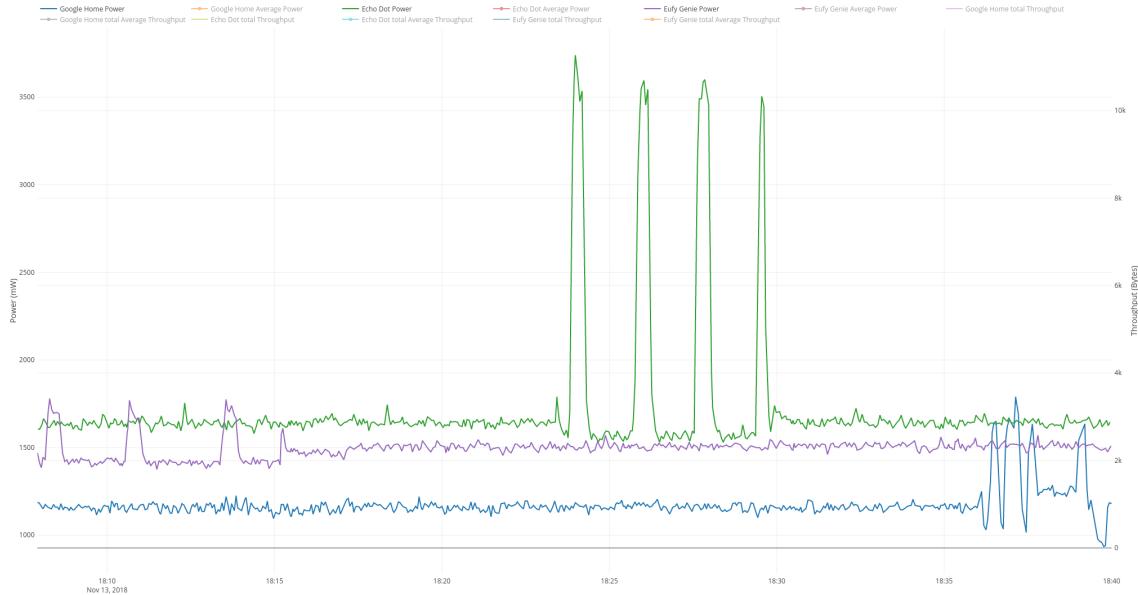


Figure 3.25: Power usage of Echo Dot 1, Eufy Genie 1, and Google Home over time when asked “who is the best basketball player”. Same graph as figure 3.15 with Echo DOT 2 and Eufy Genie removed.

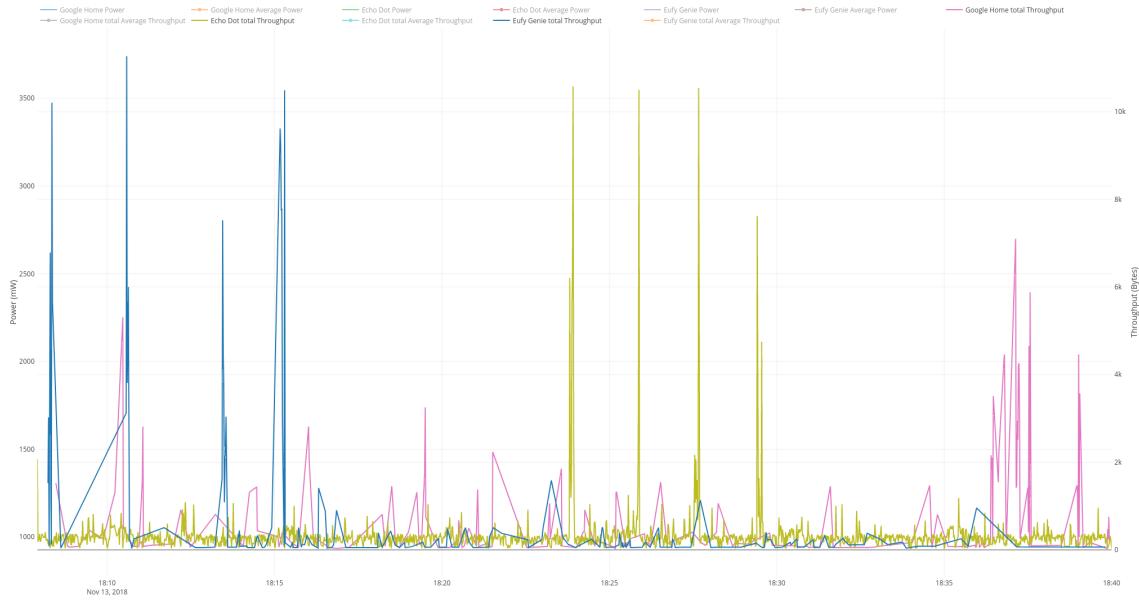


Figure 3.26: Power usage of Echo Dot 1, Eufy Genie 1, and Google Home over time when asked “who is the best basketball player”. Same time frame as figure 3.25

Chapter 4

DISCUSSION

In this chapter, we discuss the implications of the data presented in chapter 3.

4.1 Switching Smart Plug

In section 3.1, we displayed figure 3.1. We do this to ensure consistency between measurements. The network traffic is assumed to be consistent because all network traffic goes through our NUC AP, following a standard WiFi connection. However, each device has their power plug in different locations on different power strips. We wanted to check to see that this variation causes minimal recording differences. From the figure 3.1, we can see that after restarting, the power level goes back to where it was before. We did this multiple times with other switches to different locations and obtained nominal changes in the idle power usage. Thus we can conclude that this power logging scheme is relatively consistent.

4.2 Metadata Analysis

In sections 3.2, 3.3 and, 3.4 we go over the protocol spread, and general power/network usage of varying devices. However, because Frawley's paper [8] already analyzes these sections, refer to his paper for analysis.

4.3 Power Spike at 18:12

In figure 3.14, we saw an unaccounted power spike at 18:12. When looking at figures 3.15 and 3.16, we believe that the power surge is from the LEDs. Because there is no network usage in this time, we do not think the Echo Dot was doing anything data recording, listening, or was activated. We believe that the computer was covering the

Echo Dot, but is briefly moved, exposing the Echo Dot to more light, thus causing the LEDs to brighten.

4.4 Device Fingerprinting

We want to see if it is possible to determine what devices are in use if given access to home's power line. Classifying devices from a single power usage graph over time consisting of multiple devices is a good simulation for a total powerline.

In sections 3.5 and 3.6. We show how the smart speakers perform under various operations such as "what's the weather", "who is the best basketball player", or "what's the news". We did this various times to confirm consistency, and we always found that there is a start spike (voltage spike peak to peak value) of around 390 mW with 36.1 mW standard deviation for the Eufy Genie, 606.7 mW spike with 110 mW standard deviation for the Google Home, and 2026.7 mW spike with 149.89 mW standard deviation for the Echo Dot as shown in figure 3.17. Whether summing the power usage graphs for the devices up or keeping separate, from the graph, we can tell what device model is being used if it is within those thresholds. The models with multiple units are consistent with each other. Ordering the queries results in the same power spike. Because of all of this, we conclude that we can classify a device model from just a visual examination of its power usage. This is our first step to testing our hypotheses that it is possible to see what devices are in use from analysis of someone's power line.

The next step is to add noise into our graphs, to see if the smart speaker spikes during operation are still be legible 3.6. From the graphs shown, if the power is stable enough, then it is not a problem for classifying the devices. The device spikes are just shifted up. This statement is supported by our graphs with the fan, the idle microwave, idle fridge, and idle NUC.

However, if the power usage has noise on the same magnitude or larger than the smart speaker spikes, then it becomes difficult to make out which smart speaker model is currently in use. This is evident when the NUC, fridge, and microwave are in use. However, the microwave and fridge are idle for most the time, so it is still possible to

monitor someone’s power line on a time frame that these devices are not being used. However, PCs, smartphones, and streaming devices are constantly in use. It would be difficult to find a time one the NUC is not in use and make it difficult to determine what smart speaker model is in use form just visual analysis of its power usage. We think that it still might be possible with machine learning, or if network data is added in as well, or both.

From the results of section 3.5 and 3.6, it is possible to classify the model of a smart speaker from visual examination of a single power usage graph during a time when most other devices are idle. However, with multi-use devices such as a PC, smartphone, or IoT device that is in constant use for long periods of times, it would be challenging to classify a smart speaker model.

4.5 Smart Speaker Comparison

Finally, we noticed that the Echo Dot is characterized by a power spike that is a lot larger than the other smart speakers and was interested to see what trade-offs these devices made.

From the results in section 3.7, we can see that the echo uses the most power, then the Eufy, then the Google Home. However, the Eufy uses the most network throughput, than then Google Home, then the Echo Dot. We believe the Eufy is constantly communicating through the network so that it can cache results, minimizing onboard operations. The Echo Dot, on the other hand, uses the least network throughput and opts to do as much on board work as possible. The Google Home sits in the middle, querying the network, for more relevant things that it can cache for when the user asks the questions. It seems they pre-request results better. It is interesting that the Google Home still sits in the middle for network throughput when it is the only one that sends UPnP at such a high frequency. However, this is mostly speculation and would be interesting to see more research on this.

Chapter 5

CONCLUSION

To conclude, this paper shows that it is easy to tell which smart speaker is in use based on just the initial spike. Devices are each making their unique trade-off between power and network throughout when idle and in use. This tradeoff leads to unique signatures that allow us to classify devices. For the smart speakers, it is as obvious as the first power spike during usage.

Some patterns can be difficult to make out with so much noise coming from within a house for power data. However, if enough devices are idle, it is still possible. It shows an indirect security flaw in our power-line and how that defense should be fortified. This paper showed it is possible to determine what smart speaker is in use and whether the lights are on in the house from the power data. With more research, it is likely that other private information can be inferred from the power line data.

This proves that more research should be done to analyze and classify IoT devices. We created an expansive database of network and power data for 16 IoT devices that spread over ten months, accumulating 184.94 GB of data and 172,445,929 entries in the database. Along with that, we created an analysis tool to examine the database visually.

5.1 Future Work

With an ever-growing database, its maintainence is necessary. We want to clean up the database by indexing all columns for quick lookup, as of now, only some columns are indexed. There are also some inconsistencies in the database because a device was inconsistently connected to a smart plug or if the IP address changes. Changing all those values would be apart of maintenance of this database.

The next plan, is to improve the realtimeIoTgrapher and add features. Hosting it on a public domain would make this tool more easily accessible, rather than running

it locally and going to a local domain. Including more automated annotation on the graphs would also make anomaly detection easier.

We want to also utilize both power and network data in tandem for analysis to see if that yields better classification results.

With many botnet attacks in the recent past, we also attempted to hack into these devices with Mirai and analyze power usage and network throughput of these devices. Getting Mirai working in the restricted time turned out to be unfeasible, but would be interesting to continue for future work.

Finally, applying machine learning to any of the ideas above or this paper's project would be open to future work as well. We quickly tried a neural net with 0 percent accuracy and an LSTM and achieved 30 percent accuracy given five devices (slightly better than random). We had also planned to alter a human activity recognition machine learning (HAR) model to fit the database for classification of device model given power data over a set period. However, we did not have time for that, and that would be interesting for future work as well.

Regardless, there are many opened opportunities for future work because of the database. They can quickly utilize the realtimeIoTGrapher to analyze the database or manually query data. There have already been students who have used the database for various research projects.

BIBLIOGRAPHY

- [1] “IoT: Number of Connected Devices Worldwide 2012-2025,” 2016. [Online]. Available: <https://www.statista.com/statistics/471264/iot-number-of-connected-devices-worldwide/>
- [2] iotforall, “The 5 Worst Examples of IoT Hacking and Vulnerabilities in Recorded History,” 2017. [Online]. Available: <https://www.iotforall.com/5-worst-iot-hacking-vulnerabilities/>
- [3] S. Burke, “Google Admits Its New Smart Speaker Was Eavesdropping on Users,” 2017. [Online]. Available: <https://money.cnn.com/2017/10/11/technology/google-home-mini-security-flaw/index.html>
- [4] Z. Whittaker, “Fear the Reaper? Experts Reassess the Botnet’s Size and Firepower,” 2017. [Online]. Available: <https://www.zdnet.com/article/reaper-botnet-experts-reassess-size-and-firepower/>
- [5] Y. Amar, H. Haddadi, R. Mortier, A. Brown, J. Colley, and A. Crabtree, “An Analysis of Home IoT Network Traffic and Behaviour,” *ArXiv e-prints*, 2018.
- [6] Y. Meidan, M. Bohadana, A. Shabtai, J. D. Guarnizo, M. Ochoa, N. O. Tippenhauer, and Y. Elovici, “ProfilIoT: a Machine Learning Approach for IoT Device Identification Based on Network Traffic Analysis,” in *Proceedings of the Symposium on Applied Computing*, ser. SAC ’17. ACM, 2017, pp. 506–509. [Online]. Available: <http://doi.acm.org/10.1145/3019612.3019878>
- [7] Wireshark, “Wireshark.” [Online]. Available: <https://www.wireshark.org/>
- [8] R. J. Frawley, “Logging and Analysis of Internet of Things (IoT) Device Network Traffic and Power Consumption,” 2018. [Online]. Available: <https://digitalcommons.calpoly.edu/theses/1911/>
- [9] “Geoip databases and services: Industry leading ip intelligence.” [Online]. Available: <https://www.maxmind.com/en/geoip2-services-and-databases>
- [10] Amazon, “Amazon rds.” [Online]. Available: <https://aws.amazon.com/rds/>

- [11] MySQL, “Mysql.” [Online]. Available: <https://www.mysql.com/>
- [12] Intel, “Intel nuc.” [Online]. Available: <https://www.intel.com/content/www/us/en/products/boards-kits/nuc.html>
- [13] Ubuntu, “Ubuntu.” [Online]. Available: <https://www.ubuntu.com/>
- [14] linux, “linux.” [Online]. Available: <https://www.linux.org/>
- [15] Ookla, “Speedtest.” [Online]. Available: <https://www.speedtest.net/>
- [16] Google, “Chromebook.” [Online]. Available: <https://www.google.com/chromebook/>
- [17] Belkin, “Wemo.” [Online]. Available: <https://www.belkin.com/us/Products/smarthome-iot/c/wemo/>
- [18] S. Perez, “Voice-enabled Smart Speakers to Reach 55 Percent of U.S. Households by 2022, Says Report,” 2017. [Online]. Available: <https://techcrunch.com/2017/11/08/voice-enabled-smart-speakers-to-reach-55-of-u-s-households-by-2022-says-report/>
- [19] ——, “39 Million Americans Now Own a Smart Speaker, Report Claims,” 2018. [Online]. Available: <https://techcrunch.com/2018/01/12/39-million-americans-now-own-a-smart-speaker-report-claims/>
- [20] J. Lynch, “Nearly 70 Million U.S. Households Now Have a Connected-TV Streaming Device,” 2017. [Online]. Available: <https://www.adweek.com/tv-video/nearly-70-million-u-s-households-now-have-a-connected-tv-streaming-device/>
- [21] “168 Million Will Watch Connected TV in the US This Year,” 2017. [Online]. Available: <https://www.emarketer.com/Article/168-Million-Will-Watch-Connected-TV-US-This-Year/1016233>
- [22] https://www.nintendo.co.jp/ir/en/finance/hard_soft/index.html, “Dedicated Video Game Sales Units.”

- [23] A. Souppouris, “Sony Has Sold 50 Million PlayStation 4s,” 2016. [Online]. Available:
<https://www.engadget.com/2016/12/07/sony-playstation4-50-million-sales/>
- [24] “The 5 Worst Examples of IoT Hacking and Vulnerabilities in Recorded History,” 2018. [Online]. Available:
<https://www.iotforall.com/5-worst-iot-hacking-vulnerabilities/>
- [25] N. Feamster, “Who Will Secure the Internet of Things,” 2016. [Online]. Available:
Available: <https://freedom-to-tinker.com/2016/01/19/who-will-secure-the-internet-of-things/>
- [26] Oblique, “Create ap,” 2017. [Online]. Available:
https://github.com/oblique/create_ap
- [27] J. M, “How to create wifi hotspot in ubuntu 16.04,” 2016. [Online]. Available:
<http://ubuntuhandbook.org/index.php/2016/04/create-wifi-hotspot-ubuntu-16-04-android-supported/>
- [28] mysqlclient, “mysqlclient.” [Online]. Available:
<https://pypi.org/project/mysqlclient/>
- [29] pavoni, “pywemo,” 2018. [Online]. Available: <https://github.com/pavoni/pywemo>
- [30] navicat, “Navicat,” 2018. [Online]. Available: <https://www.navicat.com/en/>
- [31] mysql, “Mysql commandline,” 2018. [Online]. Available:
<https://dev.mysql.com/doc/refman/5.5/en/mysql.html>
- [32] google, “Google sheets,” 2018. [Online]. Available:
<https://www.google.com/sheets/about/>
- [33] pandas, “Pandas,” 2018. [Online]. Available: <https://pandas.pydata.org>

APPENDICES

Appendix A

HOW TO INSTALL AND RUN REALTIMEIOTGRAPH.PY

This chapter discusses how to install, run, and use the database visualization tool.

A.1 How To Install and Run

To install and run this analysis tool, follow the instructions at the GitHub site: github.com/nealhnguyen/realtimeiotgrapher.

First, pull the code for this tool from git hub and install the dependencies in the requirements file through pip. At this point, the dependencies are set to run this code on python 2.7.

Next, a 'loginCredentials' file must be created containing the host, user, password, and database information. At which point, running 'python realTimeIoTGraph.py' starts the backend and display its current operation onto the terminal.

The link to the local IP address contains the graphing tool. Clicking on that opens it onto the default browser for the computer.

A.2 Usage Directions

Once the software is running, the grapher can be viewed on a web browser as shown in figure A.1. This section describes each numbered element in the image.



Figure A.1: RealTimeIoTGrapher usage image

0. The local IP address presented in the terminal where the grapher is hosted.
1. The interval field and the text box next to it is dark when the software is in live update mode.
2. The text box here relates to the live update feature of the graph. This is a time field, denoting the window of time the graph displays at a time while it updates. The format is HH:MM:SS where HH is the hour, MM is minutes, and SS is seconds.
3. The time range field and the four elements consecutively after it are dark instead of grey when it is in static graph mode. In this state, the graphing tool specifies the network traffic and power information in the time window specified.
4. In these three text boxes, the time window for static graphing is specified. The first box is the start time, the second box is the end time, and the last box is the interval. To specify the time range, fill in any 2 of the three text boxes, and the graphing tool infers the other number. If all three text boxes are filled, then the

graphing tool uses the start and end time text box. To further explain, if the start time and interval are specified, then the window is start time → start time + interval. If the end time and interval are specified, then the time window is the end time - interval → end time.

5. This button causes the graph to update. After updating the time range in step 4, clicking this button updates the graph with the new time range.
6. The “Use Time Range” checkbox toggles the software between live graph mode and static mode. The “Sum Graph” checkbox sums up all power graphs into one. The “Show Only Power” checkbox allows the grapher to avoid a SQL query for the network throughput, speeding up the process in case of pure power analysis.
7. This selection field specifies all the devices the software creates graphs. One, or many devices can be specified at a time. If too many devices are specified at a time, the software may run slow.
8. The traces shown can be selected or deflected to toggle its visibility.
9. This line graph displays the number of bytes sent through the network by a device going in, out, and in total. It also shows the power usage over time. Also, it shows the average value for a graph in the interval of the graph
10. This bar graph displays the spread of network traffic by protocol for each device. It displays the outgoing packets as orange and incoming packets as blue. It displays an individual bar graph for each device and protocol. Similarly, the traces can be clicked on to hide its visibility.